

Ganesh Gupta

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

EXPERIENCE

Oracle | *Software Engineer, AI Platform*

Jun 2025 – Present

- Provisioned and operated multi-cloud **GPU clusters** (AWS, GCP, CoreWeave) with **NVIDIA/CUDA** and **InfiniBand**, running distributed training on **Slurm** and **Kubernetes**.
- Built observability and autoscaling for **vLLM** inference serving with **Prometheus/Grafana**, and trained and fine-tuned **PyTorch** models on multi-GPU clusters.
- Built a **distributed executor** exposing typed **REST APIs** over Docker and serverless compute, backed by an **S3** artifact store at **sub-200ms p99**, with **OAuth/RBAC**.
- Led a production **multi-agent system** on **OpenAI Agents SDK**, **LangGraph**, and **MCP** tool use that autonomously processes **financial workflows**, handling **millions of daily transactions** at **99%+ uptime**.

QureAI | *ML Infrastructure Engineer*

Mar 2026 – Present

- Extended **vLLM-class** serving internals with **CUDA/Triton** kernels, distributed **KV cache**, and **speculative decoding**, alongside **TensorRT-LLM**.
- Cut **p99 latency from 420ms to 140ms** across **50K+ daily requests** by moving LLM serving to an **async, autoscaled** architecture.
- Built **PySpark** jobs and **Airflow** DAGs for large-scale feature and dataset preparation feeding training and **evals**.

Nomura Research Institute Financial Technologies | *Software Engineer*

Sep 2019 – Aug 2023

- Architected and operated **25+ production Python and Java microservices** for **1M+ users** in **financial services** at **99.9% uptime** on **AWS**.
- Cut production escalations **60%** in one quarter by standardizing API contracts and failure handling, driving **root cause analysis** and post-incident reviews.
- Cut nightly batch aggregation from **8 minutes to 90 seconds** by optimizing **SQL and PL/SQL** on a financial RDBMS with indexed views and partitioning.
- Led **architecture and code reviews** and mentored engineers on system design, reliability, and secure API practices.

PROJECTS

Distributed KV Store

github.com/ganeshhgupta/distributed-kv-store

- **Raft-consensus** key-value store with **linearizable reads** and **log compaction**, plus a companion VPC control plane simulator.

Watch-your-LLM

github.com/ganeshhgupta/watch-your-llm

- **LLM evaluation** platform with rubric-based scoring, **hallucination detection**, regression alerting, a **FastAPI** backend, and a **React/TypeScript** dashboard.

RootCause AI

github.com/ganeshhgupta/rootcause-ai

- **LangGraph** multi-agent system that routes production alerts to specialized sub-agents via **tool-use** chains and writes structured root-cause reports, cutting **MTTR 65%**.

EDUCATION

University of Texas at Arlington | *M.S. Computer Science, GPA 4.0/4.0*

Aug 2023 – May 2025

MCKV Institute of Engineering | *B.Tech Computer Science*

Aug 2016 – Jul 2020

SKILLS

Languages: Python, C++, Go, SQL, TypeScript

Inference & ML: vLLM, TensorRT-LLM, CUDA, Triton, KV Cache, Speculative Decoding, PyTorch, Fine-Tuning, Evals

Infra: Kubernetes, Slurm, Docker, NVIDIA GPUs, InfiniBand, AWS, GCP, CoreWeave, Prometheus/Grafana, CI/CD

Data: PySpark, Airflow, PostgreSQL, Redis, Kafka, FastAPI, REST APIs

Certifications: AWS Certified Developer Associate, GCP Professional MLOps Engineer