

GANESH GUPTA

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

Experience

Oracle

Jun 2025 – Present

Software Engineer, AI Platform

- Build the agentic harness layer powering production AI workflows: task planning, tool routing, and stateful agent memory via Redis checkpointing across LangGraph, CrewAI, AutoGen, and OpenAI Agents SDK, at 99%+ uptime.
- Hands-on with Anthropic Claude APIs (tool use, structured outputs, prompt caching, extended thinking) alongside OpenAI models, selecting model/mode per agentic workflow; use Claude Code and Codex daily as core dev tools.
- Built a deployment-gated evaluation framework with guardrails, traceability, and production observability across agent outputs, blocking 15% of failed releases before they reached users.
- Designed scalable platform APIs (Python, TypeScript/Node.js) enabling product teams to build and operate agentic applications safely, backed by PostgreSQL and Redis.
- Set engineering standards for CI/CD, containerized deployment, and secure development practices across the platform, working in Agile sprints with Git-based version control.

QureAI (AI Startup)

Mar 2026 – Present

ML Infrastructure Engineer

- Built custom evaluators for RAG/agent workflows tracking drift, latency, and regression across model versions, catching failures before they reached production.
- Improved semantic retrieval recall 61% to 89% on a 2K-query eval set by rebuilding the chunking pipeline with hybrid dense-sparse retrieval and re-ranking.

Nomura Research Institute Financial Technologies

Sep 2019 – Aug 2023

Software Engineer

- Architected and operated 25+ production C++ and SQL-backed microservices for 1M+ users, sustaining 99.9% uptime over 4 years on AWS (EC2, RDS, SQS).
- Drove a redesign of nightly data aggregation, cutting processing from 8 minutes to 90 seconds via indexed views and batch partitioning.
- Standardized failure-handling contracts (retry, timeout, propagation semantics) across 25+ services, reducing production escalations 60% in one quarter.
- Mentored engineers through code reviews and design write-ups, driving consistent engineering standards.

Projects

LangFetch | tinyurl.com/LangFetch

- LangGraph agent with MCP-compatible tool-use compiling natural-language questions into dependency-ordered SQL subqueries with Pinecone vector retrieval and hybrid re-ranking; deployed as an MCP server for enterprise data queries.

RootCause AI | github.com/ganeshhgupta/rootcause-ai

- LangGraph multi-agent system diagnosing production incidents via tool-use chains; routes alerts to specialized sub-agents and synthesizes evidence-backed root cause reports, reducing mean time to resolution.

Watch-your-LLM | github.com/ganeshhgupta/watch-your-llm

- Production LLM eval platform with hallucination detection, drift tracking, rubric-based scoring, and regression alerting across model versions; FastAPI backend, PostgreSQL.

Education

University of Texas at Arlington

Aug 2023 – May 2025

MS, Computer Science, GPA 4.0/4.0

MCKV Institute of Engineering

Jul 2016 – Jul 2020

BTech, Computer Science

Skills

Languages: Python, TypeScript, SQL, Java, C++

Agentic AI & Harness: LangGraph, CrewAI, AutoGen, OpenAI Agents SDK, Anthropic Claude (tool use, structured outputs, prompt caching), MCP, Agent Memory/State, Evaluation Loops

Retrieval & Safety: RAG, GraphRAG, Semantic Retrieval, Guardrails, Model/Agent Observability, Traceability

Full-Stack & Cloud: React, Node.js, REST APIs, PostgreSQL, Redis, AWS, Azure, GCP, Docker, Kubernetes, CI/CD

Certifications

AWS Certified Developer Associate | GCP Professional MLOps Engineer