

Ganesh Gupta

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

EXPERIENCE

Oracle | *Software Engineer, AI Platform*

Jun 2025 – Present

- Owned full-stack delivery of user-facing product features end-to-end, from React/TypeScript frontend through Python/FastAPI backend services and APIs, shipping directly to production.
- Used AI coding tools (Claude Code, Codex) daily to accelerate implementation, debugging, and code review across the stack, cutting feature turnaround time.
- Replaced polling with SSE for live product metrics streaming, cutting redundant API calls 90% and improving perceived responsiveness for end users.
- Built a multi-agent system via OpenAI Agents SDK and MCP tool-use handling millions of daily user-facing transactions at 99%+ uptime.
- Debugged and resolved production issues based on user impact, partnering with product and engineering to prioritize fixes that moved the needle fastest.

QureAI | *ML Infrastructure Engineer*

Mar 2026 – Present

- Built a model observability framework tracking drift, latency, and eval regressions in production, enabling governance-aligned rollback workflows.
- Cut p99 inference latency 420ms to 140ms across 50K+ daily requests by migrating synchronous serving to async vLLM on GCP Cloud Run with autoscaling.

Nomura Research Institute Financial Technologies | *Software Engineer*

Sep 2019 – Aug 2023

- Built and operated 25+ production microservices in Java, C++, and SQL-backed distributed systems for 1M+ users, sustaining 99.9% uptime over 3 years on AWS (EC2, RDS, SQS).
- Cut nightly financial data aggregation blocking downstream settlement from 8 minutes to 90 seconds by redesigning SQL with indexed views and batch partitioning.
- Standardized failure-handling contracts across 25+ financial services, reducing production escalations 60% in one quarter via consistent retry, timeout, and propagation semantics.

PROJECTS

Watch-your-LLM

github.com/ganeshhgupta/watch-your-llm

- Full-stack LLM eval and observability product: FastAPI backend, React/TypeScript dashboard, PostgreSQL, with rubric-based scoring and drift detection shipped end-to-end solo.

RootCause AI

github.com/ganeshhgupta/rootcause-ai

- LangGraph multi-agent system automating production incident diagnosis via tool-use chains; routes alerts to specialized sub-agents and synthesizes root cause reports, reducing manual toil.

LangFetch

tinyurl.com/LangFetch

- LangGraph agent decomposing multi-table questions into dependency-ordered SQL subqueries with intermediate result passing; deployable as an MCP-compatible server for enterprise data queries.

EDUCATION

University of Texas at Arlington | *M.S. Computer Science, GPA 4.0/4.0*

Aug 2023 – May 2025

MCKV Institute of Engineering | *B.Tech Computer Science*

Jul 2016 – Jul 2020

SKILLS

Languages: Python, TypeScript, JavaScript, Java, C++, SQL

Full-Stack: React, FastAPI, Node.js, REST APIs, SSE

AI/Product: AI Coding Tools (Claude Code, Codex), LLMs (OpenAI, Anthropic), Agentic AI, RAG, LangChain/LangGraph

Data & Cloud: PostgreSQL, MongoDB, AWS (EC2, S3, SageMaker), GCP (Cloud Run), Docker, Kubernetes, CI/CD

CERTIFICATIONS

AWS Certified Developer Associate | Google Cloud Professional Machine Learning Engineer