

Ganesh Gupta

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

EXPERIENCE

Oracle | *Software Engineer, AI Platform*

Jun 2025 – Present

- Built a deployment-gated evaluation and developer-tooling platform used across the engineering team, with rubric-based benchmarking that blocked 15% of regressions before release.
- Designed and optimized agentic AI workflows across multiple LLMs (GPT, Claude, Llama, Gemini), building AI-native capabilities that help engineers plan and execute changes faster.
- Built a multi-agent system via OpenAI Agents SDK and MCP tool-use handling millions of daily transactions at 99%+ uptime.
- Engineered a pluggable distributed executor exposing typed REST APIs, routing pipeline nodes across Docker and AWS Lambda, backed by an S3 artifact store eliminating redundant recomputation.
- Drove design and code reviews enforcing style and testability standards, and triaged production issues affecting service reliability.

QureAI | *ML Infrastructure Engineer*

Mar 2026 – Present

- Built a model observability framework tracking drift, latency, and eval regressions in production, enabling governance-aligned rollback workflows.
- Cut p99 inference latency 420ms to 140ms across 50K+ daily requests by migrating synchronous serving to async vLLM on GCP Cloud Run with autoscaling.

Nomura Research Institute Financial Technologies | *Software Engineer*

Sep 2019 – Aug 2023

- Built and operated 25+ production microservices in Java, C++, and SQL-backed distributed systems for 1M+ users, sustaining 99.9% uptime over 3 years on AWS (EC2, RDS, SQS).
- Standardized failure-handling contracts across 25+ financial services, reducing production escalations 60% in one quarter via consistent retry, timeout, and propagation semantics.
- Cut nightly financial data aggregation blocking downstream settlement from 8 minutes to 90 seconds by redesigning SQL with indexed views and batch partitioning.

PROJECTS

RootCause AI

github.com/ganeshhgupta/rootcause-ai

- LangGraph multi-agent system automating production incident diagnosis via tool-use chains; routes alerts to specialized sub-agents and synthesizes root cause reports, reducing manual toil for engineers.

Watch-your-LLM

github.com/ganeshhgupta/watch-your-llm

- Production eval platform with rubric-based scoring, drift detection, and latency tracking across model versions; FastAPI backend, React/TypeScript dashboard, PostgreSQL.

LangFetch

tinyurl.com/LangFetch

- LangGraph agent decomposing multi-table questions into dependency-ordered SQL subqueries with intermediate result passing; deployable as an MCP-compatible server for enterprise data queries.

EDUCATION

University of Texas at Arlington | *M.S. Computer Science, GPA 4.0/4.0*

Aug 2023 – May 2025

MCKV Institute of Engineering | *B.Tech Computer Science*

Jul 2016 – Jul 2020

SKILLS

Languages: Java, C++, Python, TypeScript, SQL

AI/ML: LLMs (GPT, Claude, Llama, Gemini), Agentic AI, LangChain, LangGraph, OpenAI Agents SDK, MCP, Developer Tooling

Systems: Distributed Systems, Object-Oriented Design, Fault Tolerance

Cloud & Infra: AWS (EC2, RDS, SQS, Lambda, S3), GCP (Cloud Run), Docker, Kubernetes, REST APIs

CERTIFICATIONS

AWS Certified Developer Associate | Google Cloud Professional Machine Learning Engineer