

Ganesh Gupta

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

EXPERIENCE

Oracle | *Software Engineer, AI Platform*

Jun 2025 – Present

- Built and extended agentic AI workflows that reason over context, call tools, and perform actions, via LangGraph, CrewAI, AutoGen, and the OpenAI Agents SDK with MCP tool-use, at 99%+ uptime.
- Deployed a RAG pipeline on AWS SageMaker for real-time decision support at sub-200ms p99 latency with RTO under 2 minutes, with an emphasis on retrieval correctness and safety.
- Contributed to shared AI infrastructure: a deployment-gated evaluation and monitoring framework enforcing guardrails and quality/regression checks, blocking 15% of failed releases before production.
- Operated AI systems in production, monitoring and debugging a distributed event-driven executor (Docker, AWS Lambda, Kafka-based messaging) backed by an S3 artifact store.
- Used AI-assisted and agentic development tools (Claude Code, GitHub Copilot) daily for design, implementation, testing, and refactoring, while owning production-quality standards.

QureAI | *ML Infrastructure Engineer*

Mar 2026 – Present

- Built evaluation and monitoring tooling measuring LLM/agent output quality, drift, and regression across versions, catching failures before they reached production.
- Improved RAG retrieval recall 61% to 89% on a 2K-query eval set by rebuilding the chunking pipeline with hybrid dense-sparse retrieval over a vector store and re-ranking.

Nomura Research Institute Financial Technologies | *Software Engineer*

Sep 2019 – Aug 2023

- Built and operated 25+ production Java and C++ microservices for financial transaction processing, serving 1M+ users at 99.9% uptime over 3 years on AWS (EC2, RDS, SQS), in a regulated financial environment.
- Standardized failure-handling contracts across 25+ financial services, reducing production escalations 60% in one quarter via consistent retry, timeout, and propagation semantics.

PROJECTS

LangFetch

tinyurl.com/LangFetch

- LangGraph agent decomposing multi-table questions into dependency-ordered SQL subqueries with Pinecone vector retrieval and hybrid re-ranking; deployed as an MCP-compatible API server for enterprise data queries.

RootCause AI

github.com/ganeshhgupta/rootcause-ai

- LangGraph multi-agent system automating production incident diagnosis via tool-use chains; routes alerts to specialized sub-agents and synthesizes evidence-backed root cause reports.

Watch-your-LLM

github.com/ganeshhgupta/watch-your-llm

- Production LLM evaluation and monitoring platform with hallucination detection, drift tracking, rubric-based scoring, and schema-validated structured outputs; FastAPI backend, PostgreSQL.

EDUCATION

University of Texas at Arlington | *M.S. Computer Science, GPA 4.0/4.0*

Aug 2023 – May 2025

MCKV Institute of Engineering | *B.Tech Computer Science*

Jul 2016 – Jul 2020

SKILLS

Languages: Python, TypeScript, Java, C++, SQL

Agentic AI & LLMs: ChatGPT, Claude, Claude Code, Gemini, GitHub Copilot, LangChain, LangGraph, CrewAI, AutoGen, OpenAI Agents SDK, MCP, RAG, Prompt Engineering, vLLM

APIs & Data: REST APIs, Vector Storage (Pinecone), Schema Validation, PostgreSQL

Distributed Systems & Cloud: Kafka, Event-Driven Architecture, Kubernetes, Docker, AWS (EC2, S3, RDS, SQS, SageMaker), CI/CD

CERTIFICATIONS

AWS Certified Developer Associate | Google Cloud Professional Machine Learning Engineer