

Ganesh Gupta

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

EXPERIENCE

Oracle | *Software Engineer, AI Platform*

Jun 2025 – Present

- Designed and deployed agentic AI and GenAI/LLM solutions across multiple models (GPT, Claude, Llama, Gemini), selecting and tuning models per use case for production quality and cost.
- Built a multi-agent system via OpenAI Agents SDK and MCP tool-use handling millions of daily transactions at 99%+ uptime.
- Deployed a RAG pipeline on AWS SageMaker for real-time decision support at sub-200ms p99 latency with RTO under 2 minutes.
- Built a deployment-gated evaluation and MLOps framework enforcing guardrails and quality checks across model outputs, blocking 15% of failed releases before production.

QureAI | *ML Infrastructure Engineer*

Mar 2026 – Present

- Developed computer vision and multi-modal models (PyTorch) for medical image understanding, running model training and fine-tuning on multi-GPU clusters with compression and quantization.
- Improved RAG retrieval recall 61% to 89% on a 2K-query eval set by rebuilding the data pipeline with hybrid dense-sparse retrieval and re-ranking.

Nomura Research Institute Financial Technologies | *Software Engineer*

Sep 2019 – Aug 2023

- Built and operated 25+ production microservices in Java, C++, and SQL-backed distributed systems for 1M+ users, sustaining 99.9% uptime over 3 years on AWS (EC2, RDS, SQS).
- Cut nightly financial data aggregation blocking downstream settlement from 8 minutes to 90 seconds by redesigning SQL with indexed views and batch partitioning.

PROJECTS

LLaMA 2 Fine-Tuning Pipeline for Financial Documents

github.com/ganeshhgupta/llama2-finance-finetuning

- Post-trained LLaMA 2 with LoRA and PEFT for domain-specific financial document understanding, reducing trainable parameter count while preserving model capacity via systematic hyperparameter optimization.

Watch-your-LLM

github.com/ganeshhgupta/watch-your-llm

- Production LLM eval and observability platform with rubric-based scoring, drift detection, and latency tracking across model versions; FastAPI backend, PostgreSQL.

LangFetch

tinyurl.com/LangFetch

- LangGraph agent decomposing multi-table questions into dependency-ordered SQL subqueries with intermediate result passing; deployed as an MCP-compatible API server for enterprise data queries.

EDUCATION

University of Texas at Arlington | *M.S. Computer Science, GPA 4.0/4.0*

Aug 2023 – May 2025

MCKV Institute of Engineering | *B.Tech Computer Science*

Jul 2016 – Jul 2020

SKILLS

Languages: Python, Java, TypeScript, C++, SQL

AI/ML: LLMs, GenAI, Agentic AI, PyTorch, RAG, LoRA/PEFT Fine-Tuning, Model Evaluation, Computer Vision

MLOps & Data: Model Deployment, Data Engineering, Data Pipelines, PostgreSQL, SQL/NoSQL

Cloud & DevOps: AWS (EC2, S3, SageMaker), GCP, Docker, Kubernetes, REST APIs, Git, CI/CD

CERTIFICATIONS

AWS Certified Developer Associate | Google Cloud Professional Machine Learning Engineer