

GANESH GUPTA

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

Experience

Oracle

Jun 2025 – Present

Software Engineer, AI Platform

- Lead design and delivery of multi-agent orchestration workflows (LangGraph, CrewAI, AutoGen, OpenAI Agents SDK) powering an agentic AI program used across product teams.
- Own the evaluation, guardrails, and observability layer deciding release safety: a deployment-gated framework measuring quality, latency, and regression, blocking 15% of failed releases before rollout.
- Combine GraphRAG with vector-based RAG for grounded retrieval, building reusable MCP-compatible SDK components adopted by other engineering teams.
- Drive full-stack platform delivery (React/TypeScript, Python/FastAPI), owning CI/CD, Terraform infrastructure as code, and Agile sprint execution end-to-end.

QureAI (AI Startup)

Mar 2026 – Present

ML Infrastructure Engineer

- Improved semantic retrieval recall 61% to 89% on a 2K-query eval set by rebuilding the chunking pipeline with hybrid dense-sparse retrieval and re-ranking.
- Built custom evaluators tracking drift, latency, and regression across agent/model versions, providing the trusted signal for safe rollout.

Nomura Research Institute Financial Technologies

Sep 2019 – Aug 2023

Software Engineer

- Architected and led delivery of 25+ production Java Spring Boot microservices exposing REST APIs, serving 1M+ users at 99.9% uptime on AWS (EC2, RDS, SQS).
- Standardized failure-handling contracts (retry, timeout, propagation semantics) across 25+ distributed services, reducing production escalations 60% in one quarter.
- Mentored engineers through code reviews and design write-ups, driving consistent engineering standards.

Projects

RootCause AI | github.com/ganeshhgupta/rootcause-ai

- LangGraph multi-agent system diagnosing production incidents via tool-use chains; routes alerts to specialized sub-agents and synthesizes evidence-backed root cause reports.

LangFetch | tinyurl.com/LangFetch

- LangGraph agent with MCP-compatible tool-use compiling natural-language questions into dependency-ordered SQL subqueries with vector retrieval and hybrid re-ranking; deployed as a reusable MCP server.

Watch-your-LLM | github.com/ganeshhgupta/watch-your-llm

- Production LLM eval platform with hallucination detection, drift tracking, rubric-based scoring, and regression alerting across model versions; FastAPI backend, PostgreSQL.

Education

University of Texas at Arlington

Aug 2023 – May 2025

MS, Computer Science, GPA 4.0/4.0

MCKV Institute of Engineering

Jul 2016 – Jul 2020

BTech, Computer Science

Skills

Languages: Python, TypeScript, SQL, Java, C++

Agentic AI & GenAI: LangGraph, CrewAI, AutoGen, OpenAI Agents SDK, MCP, Hybrid RAG, Vector Retrieval, Re-ranking

AI Dev Tools: Claude Code, Codex, GitHub Copilot

Full-Stack & Cloud: React, FastAPI, REST APIs, PostgreSQL, Redis, AWS, GCP, Docker, Kubernetes, Terraform, CI/CD

Certifications

AWS Certified Developer Associate | GCP Professional MLOps Engineer