

Ganesh Gupta

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

EXPERIENCE

Oracle | *Software Engineer, AI Platform*

Jun 2025 – Present

- Owned GPU fleet placement end to end: turned cluster utilization and capacity telemetry across AWS, GCP, and CoreWeave into node-pool sizing and quota decisions the platform acted on, balancing training/inference demand against fleet efficiency.
- Wrote systems-level Go and C++ scheduling logic implementing bin-packing and fault-tolerant placement heuristics under real GPU-scarcity constraints, directly improving fleet utilization across tens of thousands of jobs.
- Instrumented GPU cluster telemetry (utilization, memory, network) via Prometheus/Grafana and ran Slurm-based HPC scheduling, reasoning quantitatively about resource-consumption distributions, not just averages, to keep the fleet utilized through node loss and provider hiccups.
- Built a deployment-gated evaluation platform with staged rollout gating and rubric-based benchmarking that blocked 15% of regressions before production, with an explicit rollback path when a change misbehaved.

QureAI | *ML Infrastructure Engineer*

Mar 2026 – Present

- Owned an inference-serving change end to end – data analysis, migration to async serving with autoscaling, production rollout, and measurement – cutting p99 latency 420ms to 140ms across 50K+ daily requests on GCP Cloud Run.
- Built a model observability framework tracking drift, latency, and eval regressions in production; when a model's predictions misbehaved, diagnosed the cause, bounded the impact, and triggered governance-aligned automated rollback.

Nomura Research Institute Financial Technologies | *Software Engineer*

Sep 2019 – Aug 2023

- Built and operated 25+ production microservices for 1M+ users, sustaining 99.9% uptime over 3 years on AWS (EC2, RDS, SQS).
- Standardized failure-handling contracts across 25+ financial services, reducing production escalations 60% via consistent retry, timeout, and propagation semantics.

PROJECTS

Distributed KV Store / VPC Control Plane Simulator

github.com/ganeshhgupta/distributed-kv-store

- Raft-consensus distributed KV store handling concurrent writes with linearizable reads and log compaction; companion VPC control plane simulator modeling subnet routing and security group enforcement.

RootCause AI

github.com/ganeshhgupta/rootcause-ai

- Multi-agent system automating root-cause diagnosis of production failures via tool-use chains across distributed services, bounding impact and accelerating fixes.

Watch-your-LLM

github.com/ganeshhgupta/watch-your-llm

- Production eval and monitoring platform tracking latency, drift, and regressions across model versions and serving infrastructure; FastAPI backend, PostgreSQL.

EDUCATION

University of Texas at Arlington | *M.S. Computer Science, GPA 4.0/4.0*

Aug 2023 – May 2025

MCKV Institute of Engineering | *B.Tech Computer Science*

Jul 2016 – Jul 2020

SKILLS

Languages: Go, Python, C++, Java, SQL

Systems: Distributed Systems, OS Fundamentals, Fault Tolerance, GPU/HPC Cluster Scheduling

Capacity & Performance: Bin Packing, Node Pool/Quota Sizing, Slurm, Quantitative Resource Analysis, Prometheus, Grafana

Cloud & Infra: AWS, GCP, Docker, Kubernetes

CERTIFICATIONS

AWS Certified Developer Associate | Google Cloud Professional Machine Learning Engineer