

Ganesh Gupta

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

EXPERIENCE

Oracle | *Software Engineer, AI Platform*

Jun 2025 – Present

- Owned GPU fleet capacity end to end: translated cluster utilization and demand telemetry across AWS, GCP, and CoreWeave into node-pool sizing and quota decisions, balancing training/inference workload demand against fleet efficiency.
- Built observability and autoscaling for inference serving (Prometheus/Grafana), using metrics and logs to understand resource consumption and operational behavior across a multi-cloud fleet.
- Wrote systems-level scheduling and placement tooling (Go, C++) implementing bin-packing and fault-tolerant placement heuristics under real GPU-scarcity constraints, directly improving fleet utilization across tens of thousands of jobs.
- Applied LLM-based automation to engineering workflows (OpenAI Agents SDK, MCP tool-use), deriving insights from operational data to reduce manual triage at 99%+ uptime.

QureAI | *ML Infrastructure Engineer*

Mar 2026 – Present

- Ran Slurm-based HPC scheduling across a multi-cloud GPU fleet, keeping utilization up through node loss and provider-level disruption.
- Cut p99 serving latency 420ms to 140ms across 50K+ daily requests by diagnosing resource bottlenecks and redesigning a synchronous service into an async architecture.

Nomura Research Institute Financial Technologies | *Software Engineer*

Sep 2019 – Aug 2023

- Architected and operated 25+ production Java/C++ distributed services for 1M+ users at 99.9% uptime over 4 years on AWS (EC2, RDS, SQS), navigating complex, interdependent application architectures.
- Standardized REST API contracts and failure-handling across 25+ services, reducing production escalations 60% through repeatable engineering methodologies adopted across teams.

PROJECTS

Distributed KV Store / VPC Control Plane Simulator

github.com/ganeshhgupta/distributed-kv-store

- Raft-consensus distributed storage engine in C++ handling concurrent writes with linearizable reads and log compaction; companion VPC control plane simulator modeling subnet routing and resource allocation across nodes.

Watch-your-LLM

github.com/ganeshhgupta/watch-your-llm

- Production observability platform tracking drift, latency, and resource regressions across deployments; FastAPI backend, PostgreSQL, Docker-based deployment.

RootCause AI

github.com/ganeshhgupta/rootcause-ai

- Multi-agent system automating production incident diagnosis via tool-use chains, applying GenAI to derive root-cause insights from complex, interdependent system behavior.

EDUCATION

University of Texas at Arlington | *M.S. Computer Science, GPA 4.0/4.0*

Aug 2023 – May 2025

MCKV Institute of Engineering | *B.Tech Computer Science*

Jul 2016 – Jul 2020

SKILLS

Languages: Go, Python, Java, C++, SQL

Capacity & Infra Efficiency: Fleet Utilization/Capacity Modeling, Bin-Packing/Placement Heuristics, Demand Forecasting (GPU Infra), Multi-Cloud (AWS, GCP, CoreWeave)

Distributed Systems: Raft Consensus, Microservices, Containerized Environments (Docker, Kubernetes)

Observability: Prometheus, Grafana, Metrics/Logs-Driven Diagnosis

AI-Driven Automation: OpenAI Agents SDK, MCP, LLM-Based Insight Generation

CERTIFICATIONS

AWS Certified Developer Associate | Google Cloud Professional Machine Learning Engineer