

GANESH GUPTA

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

Experience

Oracle

Jun 2025 – Present

Software Engineer, AI Platform

- Architect distributed AI infrastructure at production scale: GPU/HPC compute clusters with InfiniBand interconnects, orchestrated via Slurm, supporting distributed workloads across cloud environments.
- Built a deployment-gated evaluation pipeline enforcing automated quality, safety, and drift checks before every release, blocking 15% of releases that failed validation, guarding system reliability.
- Drive engineering standards across the team through design reviews and code reviews, mentoring engineers on distributed system design and production reliability practices.

QureAI (AI Startup)

Mar 2026 – Present

ML Infrastructure Engineer

- Led migration of synchronous model serving to async, autoscaled architecture, cutting p99 latency 420ms to 140ms across 50K+ daily requests through distributed, high-concurrency system design.
- Diagnosed a user-facing quality gap directly from real usage signals and owned the fix end-to-end, lifting recall 61% to 89% and cutting the complaint pattern from 22% to 6% of sessions.

Nomura Research Institute Financial Technologies

Sep 2019 – Aug 2023

Software Engineer

- Architected and operated 25+ production C++ and Java microservices in a distributed, fault-tolerant system for 1M+ users, sustaining 99.9% uptime over 4 years on AWS (EC2, RDS, SQS) with relational and NoSQL data stores.
- Drove a redesign of nightly financial data aggregation across large-scale distributed data stores, cutting processing from 8 minutes to 90 seconds via indexed views and batch partitioning.
- Led standardization of failure-handling contracts across 25+ services in partnership with cross-functional teams through code reviews and design write-ups, reducing production escalations 60% in one quarter.

Projects

Distributed KV Store / VPC Control Plane Simulator | github.com/ganeshhgupta/distributed-kv-store

- Built a Raft-consensus key-value store in-house with leader election, log replication, and log compaction; companion VPC control plane simulator modeling subnet routing and security group enforcement for fault-tolerant distributed coordination.

Watch-your-LLM | github.com/ganeshhgupta/watch-your-llm

- Production LLM eval platform with hallucination detection, drift tracking, rubric-based scoring, and regression alerting across model versions; FastAPI backend, PostgreSQL, built to keep agentic systems trustworthy at volume.

RootCause AI | github.com/ganeshhgupta/rootcause-ai

- LangGraph multi-agent system diagnosing production incidents via tool-use chains; routes alerts to specialized sub-agents and synthesizes evidence-backed, traceable root cause reports.

Education

University of Texas at Arlington

Aug 2023 – May 2025

MS, Computer Science, GPA 4.0/4.0

MCKV Institute of Engineering

Jul 2016 – Jul 2020

BTech, Computer Science

Skills

Languages: C++, Java, Python, SQL

Distributed Systems: Fault Tolerance, High Availability, Consensus Protocols (Raft), Distributed Scheduling, GPU/HPC Clusters (InfiniBand, Slurm)

Backend: Java Microservices, REST APIs, PostgreSQL, NoSQL

Cloud & Infra: AWS, Azure, GCP, Docker, Kubernetes, CI/CD

Certifications

AWS Certified Developer Associate | GCP Professional MLOps Engineer