

Ganesh Gupta

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

EXPERIENCE

Oracle | *Software Engineer, AI Platform*

Jun 2025 – Present

- Built a multi-agent GenAI system via OpenAI Agents SDK and MCP tool-use, integrating ML models into production workflows handling millions of daily transactions at 99%+ uptime.
- Designed and optimized agentic AI workflows across multiple LLMs (GPT, Claude, Llama, Gemini), balancing model quality against latency and compute cost per use case.
- Deployed a production RAG service for real-time decision support at sub-200ms p99 latency with RTO under 2 minutes.

QureAI | *ML Infrastructure Engineer*

Mar 2026 – Present

- Built and deployed PyTorch-based computer vision models for medical imaging analysis, optimizing model architecture for inference efficiency in a production, resource-constrained serving environment.
- Extended and optimized inference framework internals (vLLM-class serving) with kernel-level development (CUDA/Triton), reducing compute footprint and improving throughput under latency-constrained serving.
- Ran model compression and quantization on multi-GPU clusters, cutting model size and computational cost for efficient inference, and cut p99 inference latency 420ms to 140ms across 50K+ daily requests.

Nomura Research Institute Financial Technologies | *Software Engineer*

Sep 2019 – Aug 2023

- Built and operated 25+ production microservices in Java, C++, and SQL-backed distributed systems for 1M+ users, sustaining 99.9% uptime over 3 years on AWS (EC2, RDS, SQS).
- Standardized failure-handling contracts across 25+ financial services, reducing production escalations 60% in one quarter via consistent retry, timeout, and propagation semantics.

PROJECTS

Watch-your-LLM

github.com/ganeshhgupta/watch-your-llm

- Production LLM eval platform with rubric-based scoring, drift detection, and latency tracking across model versions, designing experiments to evaluate quality/performance tradeoffs.

RootCause AI

github.com/ganeshhgupta/rootcause-ai

- Multi-agent system diagnosing production incidents via tool-use chains; routes alerts to specialized sub-agents and synthesizes root cause reports.

Distributed KV Store / VPC Control Plane Simulator

github.com/ganeshhgupta/distributed-kv-store

- Raft-consensus distributed KV store handling concurrent writes with linearizable reads and log compaction; companion VPC control plane simulator modeling subnet routing and security group enforcement.

EDUCATION

University of Texas at Arlington | *M.S. Computer Science, GPA 4.0/4.0*

Aug 2023 – May 2025

MCKV Institute of Engineering | *B.Tech Computer Science*

Jul 2016 – Jul 2020

SKILLS

Languages: Python, C++, Java, TypeScript, SQL

ML & Computer Vision: PyTorch, Computer Vision, Medical Imaging Analysis, Model Training/Fine-Tuning, LLMs

Inference Optimization: CUDA, Triton (Kernel Development), Model Compression/Quantization, Multi-GPU Training, vLLM

Systems: Distributed Systems, Operating Systems, Data Structures and Algorithms

Cloud & DevOps: AWS (EC2, S3, SageMaker), GCP (Cloud Run, Vertex AI), Docker, Kubernetes

CERTIFICATIONS

AWS Certified Developer Associate | Google Cloud Professional Machine Learning Engineer