

GANESH GUPTA

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

Experience

Oracle

Jun 2025 – Present

Software Engineer, AI Platform

- Build agentic AI systems (LangGraph, CrewAI, AutoGen, OpenAI Agents SDK) with RAG, MCP, and tool-calling integrations, deployed via Docker/CI-CD at 99%+ uptime.
- Combine GraphRAG with vector-based RAG (embeddings, semantic search) for grounded retrieval on complex enterprise queries.
- Built a deployment-gated evaluation framework enforcing guardrails and observability across model/agent outputs, blocking 15% of failed releases before production.
- Build backend services and REST APIs in Python, working in Agile sprints with Git-based version control.

QureAI (AI Startup)

Mar 2026 – Present

ML Infrastructure Engineer

- Developed deep learning models (PyTorch) with CNN-based architectures for computer vision on medical imaging data, including OCR-style document/scan analysis.
- Ran hands-on model training/fine-tuning on multi-GPU clusters with compression and quantization for efficient inference.
- Built PySpark data processing jobs for large-scale feature engineering feeding downstream model training and evaluation pipelines.

Nomura Research Institute Financial Technologies

Sep 2019 – Aug 2023

Software Engineer

- Architected and operated 25+ production C++ and SQL-backed microservices for 1M+ users, sustaining 99.9% uptime over 4 years on AWS (EC2, RDS, SQS).
- Drove a redesign of nightly data aggregation, cutting processing from 8 minutes to 90 seconds via indexed views and batch partitioning.
- Standardized failure-handling contracts (retry, timeout, propagation semantics) across 25+ services, reducing production escalations 60% in one quarter.

Projects

Watch-your-LLM | github.com/ganeshhgupta/watch-your-llm

- Production LLM eval platform with hallucination detection, drift tracking, rubric-based scoring, and regression alerting across model versions; FastAPI backend, PostgreSQL.

LangFetch | tinyurl.com/LangFetch

- LangGraph agent with MCP-compatible tool-use compiling natural-language questions into dependency-ordered SQL subqueries with Pinecone vector retrieval and hybrid re-ranking; deployed as an MCP server.

RootCause AI | github.com/ganeshhgupta/rootcause-ai

- LangGraph multi-agent system diagnosing production incidents via tool-use chains; routes alerts to specialized sub-agents and synthesizes evidence-backed root cause reports.

Education

University of Texas at Arlington

Aug 2023 – May 2025

MS, Computer Science, GPA 4.0/4.0

MCKV Institute of Engineering

Jul 2016 – Jul 2020

BTech, Computer Science

Skills

Languages: Python, SQL, TypeScript, Java, C++

ML & Deep Learning: PyTorch, scikit-learn, PySpark, Feature Engineering, CNN, Computer Vision, Model Training/Fine-Tuning

GenAI & Agentic AI: LangGraph, CrewAI, AutoGen, OpenAI Agents SDK, RAG, MCP, Prompt Engineering, Embeddings

Backend & Cloud: FastAPI, REST APIs, Docker, AWS, GCP, PostgreSQL, CI/CD

Certifications

AWS Certified Developer Associate | GCP Professional MLOps Engineer