

# GANESH GUPTA

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

## Experience

### Oracle

Jun 2025 – Present

*Software Engineer, AI Platform*

- Build multi-step autonomous agentic workflows (LangGraph, CrewAI, AutoGen, OpenAI Agents SDK) with RAG pipelines, deployed on Azure with Docker and Kubernetes.
- Integrate and orchestrate MCP (Model Context Protocol) tools connecting agents to enterprise data sources and internal APIs.
- Built a deployment-gated evaluation framework enforcing guardrails and observability across LLM/agent outputs, blocking 15% of failed releases before production.
- Turn architectural blueprints into scalable production systems (Python, FastAPI), working in Agile sprints with CI/CD.

### QureAI (AI Startup)

Mar 2026 – Present

*ML Infrastructure Engineer*

- Built custom evaluators for RAG/agent workflows tracking drift, latency, and regression across model versions, catching failures before they reached production.
- Improved semantic retrieval recall 61% to 89% on a 2K-query eval set by rebuilding the chunking pipeline with hybrid dense-sparse retrieval and re-ranking.

### Nomura Research Institute Financial Technologies

Sep 2019 – Aug 2023

*Software Engineer*

- Architected and operated 25+ production C++ and SQL-backed microservices for 1M+ users, sustaining 99.9% uptime over 4 years on AWS (EC2, RDS, SQS).
- Drove a redesign of nightly data aggregation, cutting processing from 8 minutes to 90 seconds via indexed views and batch partitioning.
- Standardized failure-handling contracts (retry, timeout, propagation semantics) across 25+ services, reducing production escalations 60% in one quarter.
- Mentored engineers through code reviews and design write-ups, driving consistent engineering standards.

## Projects

### RootCause AI | [github.com/ganeshhgupta/rootcause-ai](https://github.com/ganeshhgupta/rootcause-ai)

- LangGraph multi-agent system diagnosing production incidents via tool-use chains; routes alerts to specialized sub-agents and synthesizes evidence-backed root cause reports, reducing mean time to resolution.

### LangFetch | [tinyurl.com/LangFetch](https://tinyurl.com/LangFetch)

- LangGraph agent with MCP-compatible tool-use compiling natural-language questions into dependency-ordered SQL subqueries with Pinecone vector retrieval and hybrid re-ranking; deployed as an MCP server.

### Watch-your-LLM | [github.com/ganeshhgupta/watch-your-llm](https://github.com/ganeshhgupta/watch-your-llm)

- Production LLM eval platform with hallucination detection, drift tracking, rubric-based scoring, and regression alerting across model versions; FastAPI backend, PostgreSQL.

## Education

### University of Texas at Arlington

Aug 2023 – May 2025

*MS, Computer Science, GPA 4.0/4.0*

### MCKV Institute of Engineering

Jul 2016 – Jul 2020

*BTech, Computer Science*

## Skills

**Languages:** Python, TypeScript, SQL, Java, C++

**Agentic AI & LLMs:** LangGraph, CrewAI, AutoGen, OpenAI Agents SDK, RAG, MCP, Prompt Engineering

**Evaluation & Ops:** Model/Agent Evaluation, Guardrails, Observability, MLOps

**Cloud & DevOps:** Azure, AWS, Docker, Kubernetes, FastAPI, CI/CD

## Certifications

AWS Certified Developer Associate | GCP Professional MLOps Engineer