

# Ganesh Gupta

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

## EXPERIENCE

---

### Oracle | *Software Engineer, AI Platform*

Jun 2025 – Present

- Engineered pluggable **distributed executor** routing pipeline nodes to **Docker** or **AWS Lambda**, backed by **S3** artifact store eliminating redundant recomputation.
- Replaced polling with **SSE** for live ML metrics streaming, cutting redundant API calls **90%** and enabling **real-time** governance dashboards.
- Designed stateful recovery via per-node **Redis checkpointing**; interrupted workflows **resume from the last committed step** across **distributed execution**.
- Deployed and operated production AI workloads on **Azure** alongside **AWS**, using Azure-native services for model hosting and **pipeline orchestration**.

### QureAI | *ML Infrastructure Engineer*

Mar 2026 – Present

- Built and maintained **Airflow DAGs** orchestrating multi-stage data pipelines on **GCP**, with **dependency management** and automated **failure alerting**.
- Designed **BigQuery** schemas and optimized analytical queries for ML feature pipelines, reducing compute costs through **partition pruning** and **clustering**.

### Nomura Research Institute Financial Technologies | *Software Engineer*

Sep 2019 – Aug 2023

- Designed high-throughput SQL pipelines and **distributed messaging** workflows on **AWS SQS**, cutting nightly batch aggregation from **8 minutes to 90 seconds** via indexed views and partitioning.
- Cut nightly financial data aggregation blocking downstream **settlement** from **8 minutes to 90 seconds** by redesigning SQL with **indexed views** and **batch partitioning**.
- Standardized failure-handling contracts across **25+ financial services**, reducing production escalations **60%** in one quarter via consistent **retry, timeout, and propagation** semantics.

## PROJECTS

---

### LangFetch

[tinyurl.com/LangFetch](https://tinyurl.com/LangFetch)

- LangGraph** agent decomposes multi-table questions into dependency-ordered **SQL subqueries** with intermediate result passing; deployable as **MCP-compatible** server.

### Watch-your-LLM

[github.com/ganeshhgupta/watch-your-llm](https://github.com/ganeshhgupta/watch-your-llm)

- Production **LLM eval platform** with rubric-based scoring, **drift detection**, and latency tracking across model versions; **FastAPI**, **React/TypeScript**, **PostgreSQL**.

### RootCause AI

[github.com/ganeshhgupta/rootcause-ai](https://github.com/ganeshhgupta/rootcause-ai)

- LangGraph** multi-agent system diagnosing **production incidents** via tool-use chains; routes alerts to specialized sub-agents and synthesizes **root cause reports**, reducing **MTTR**.

## EDUCATION

---

### University of Texas at Arlington | *M.S. Computer Science, GPA 4.0/4.0*

Aug 2023 – May 2025

### MCKV Institute of Engineering | *B.Tech Computer Science*

Jul 2016 – Jul 2020

## SKILLS

---

**Languages:** Python, Java, SQL, Go, TypeScript, HTML/CSS

**Data Platform:** Airflow, BigQuery, PySpark, Kafka, Data Modeling, Data Quality and Freshness Monitoring

**Infra:** AWS, GCP, Terraform, Kubernetes, Docker, PostgreSQL, Redis

**Practices:** Data Observability, SLA Dashboards, System Design, On-call, CI/CD

## CERTIFICATIONS

---

AWS Certified Developer Associate | GCP Professional MLOps Engineer