

Ganesh Gupta

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

EXPERIENCE

Oracle | *Software Engineer, AI Platform*

Jun 2025 – Present

- Architected agentic workflows and multi-agent systems via LangGraph, CrewAI, AutoGen, and the OpenAI Agents SDK, applying supervisor/worker and hierarchical orchestration patterns with MCP tool/data access at 99%+ uptime.
- Designed context engineering strategies (context layering, compression, memory management via Redis checkpointing) to maximize reliability and token efficiency in production agent systems.
- Built a deployment-gated evaluation and observability framework enforcing guardrails, agent isolation, and quality/regression checks, blocking 15% of failed releases before production.
- Designed and optimized GenAI workflows across multiple foundation models (GPT, Claude, Llama, Gemini), architecting platform APIs (Python) enabling teams to build and launch GenAI applications end-to-end.
- Mentored engineers through code reviews and design discussions, fostering technical excellence across the agentic AI codebase.

QureAI | *ML Infrastructure Engineer*

Mar 2026 – Present

- Designed and implemented a Neo4j knowledge graph and GraphRAG pipeline modeling domain entities and relationships, enabling multi-hop reasoning and traceable, grounded retrieval in production.
- Escalated agent-to-agent task delegation via the A2A protocol, connecting a primary agent to a specialized SQL agent for grounded, explainable query execution.
- Improved RAG retrieval recall 61% to 89% on a 2K-query eval set by rebuilding the chunking pipeline with hybrid dense-sparse retrieval, multi-vector search, and re-ranking.

Nomura Research Institute Financial Technologies | *Software Engineer*

Sep 2019 – Aug 2023

- Built and operated 25+ production microservices in Java, C++, and SQL-backed distributed systems for 1M+ users in a regulated financial environment, sustaining 99.9% uptime over 3 years on AWS.

PROJECTS

LangFetch

tinyurl.com/LangFetch

- LangGraph agent with MCP-compatible tool-use compiling natural-language questions into dependency-ordered SQL subqueries with Pinecone vector retrieval and hybrid re-ranking; deployed as an MCP server for enterprise data queries.

RootCause AI

github.com/ganeshhgupta/rootcause-ai

- LangGraph multi-agent system automating production incident diagnosis via tool-use chains; routes alerts to specialized sub-agents and synthesizes evidence-backed root cause reports.

Watch-your-LLM

github.com/ganeshhgupta/watch-your-llm

- Production LLM/agent evaluation and observability platform with hallucination detection, drift tracking, rubric-based scoring, and tracing across model and agent versions; FastAPI, PostgreSQL.

EDUCATION

University of Texas at Arlington | *M.S. Computer Science, GPA 4.0/4.0*

Aug 2023 – May 2025

MCKV Institute of Engineering | *B.Tech Computer Science*

Jul 2016 – Jul 2020

SKILLS

Languages: Python, TypeScript, Java, C++, SQL

Agentic AI: LangGraph, CrewAI, AutoGen, OpenAI Agents SDK, MCP, A2A, Multi-Agent Orchestration, Context Engineering

GenAI & Retrieval: LLMs (GPT, Claude, Llama, Gemini), RAG, GraphRAG, Neo4j, LangChain, Prompt Engineering, Vector Databases, NLP

Observability & Governance: Agent/Model Evaluation, Guardrails, Agent Isolation, Deployment Gating

Cloud & DevOps: AWS, GCP, Docker, Kubernetes, CI/CD, PostgreSQL

CERTIFICATIONS

AWS Certified Developer Associate | Google Cloud Professional Machine Learning Engineer