

GANESH GUPTA

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

Experience

Oracle

Jun 2025 – Present

Software Engineer, AI Platform

- Technical anchor for the agentic AI platform, architecting multi-agent orchestration using LangGraph, LangChain, and OpenAI Agents SDK with MCP tool integration at 99%+ uptime.
- Built a deployment-gated evaluation pipeline enforcing guardrails, AI governance, and observability around agent outputs, blocking 15% of failed releases.
- Built production-grade Java and Python services with tool/function-calling integrations to internal and external systems, exposing REST APIs on multi-tenant AWS and Azure at sub-200ms latency.
- Designed stateful agent recovery via Redis checkpointing; interrupted workflows resume from the last committed step across distributed execution.
- Set engineering standards for CI/CD, testing, and security across the platform, mentoring engineers through architecture reviews.

QureAI (AI Startup)

Mar 2026 – Present

ML Infrastructure Engineer

- Built voice-based conversational AI agents on the agentic product surface, migrating synchronous model serving to async vLLM with autoscaling, cutting p99 latency 420ms to 140ms.
- Improved RAG retrieval recall 61% to 89% on a 2K-query eval set by rebuilding the chunking pipeline with hybrid dense-sparse retrieval and re-ranking.

Nomura Research Institute Financial Technologies

Sep 2019 – Aug 2023

Software Engineer

- Architected and operated 25+ production C++ and SQL-backed microservices for 1M+ users, sustaining 99.9% uptime over 4 years on AWS (EC2, RDS, SQS).
- Drove a redesign of nightly data aggregation, cutting processing from 8 minutes to 90 seconds via indexed views and batch partitioning.
- Standardized failure-handling contracts (retry, timeout, propagation semantics) across 25+ services, reducing production escalations 60% in one quarter.
- Mentored engineers through code reviews and design write-ups, driving consistent engineering standards.

Projects

LangFetch | tinyurl.com/LangFetch

- LangGraph agent compiling natural-language questions into dependency-ordered SQL subqueries with Pinecone vector retrieval and hybrid re-ranking; deployed as an MCP-compatible server for enterprise data queries.

RootCause AI | github.com/ganeshhgupta/rootcause-ai

- LangGraph multi-agent system diagnosing production incidents via tool-use chains; routes alerts to specialized sub-agents and synthesizes evidence-backed root cause reports, reducing mean time to resolution.

Watch-your-LLM | github.com/ganeshhgupta/watch-your-llm

- Production LLM eval platform with hallucination detection, drift tracking, rubric-based scoring, and regression alerting across model versions; FastAPI backend, PostgreSQL.

Education

University of Texas at Arlington

Aug 2023 – May 2025

MS, Computer Science, GPA 4.0/4.0

MCKV Institute of Engineering

Jul 2016 – Jul 2020

BTech, Computer Science

Skills

Top Skills: Python, Java, LangGraph, vLLM, Agentic AI, LLM/GenAI, Multi-Agent Systems, MCP, Conversational AI

Agentic AI & Retrieval: LangChain, CrewAI, AutoGen, RAG, Tool/Function Calling, AI Governance

Backend & Security: REST APIs, Microservices, FastAPI, Vector DBs (Pinecone, FAISS), RBAC, Guardrails

Cloud & Infra: Azure, AWS, GCP, Docker, Kubernetes, CI/CD

Certifications

AWS Certified Developer Associate | GCP Professional MLOps Engineer