

GANESH GUPTA

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

Experience

Oracle

Jun 2025 – Present

Software Engineer, AI Platform

- Lead development of agentic AI (LangGraph, OpenAI Agents SDK) test-generation and failure-analysis tooling, using prompt engineering and structured outputs to drive automated quality gates.
- Built a deployment-gated evaluation pipeline (LLM evaluation, failure clustering, root-cause reporting) enforcing guardrails and observability, blocking 15% of failed releases before production.
- Use Claude Code and Codex daily as AI-assisted development tools, accelerating test-script authoring, debugging, and PR review workflows across the platform.
- Set up GitHub Actions CI/CD pipelines, reusable workflows, and PR checks for automated testing and deployment on multi-tenant AWS and Azure infrastructure.
- Designed automated test suites driving up test coverage and release confidence, establishing engineering standards through architecture reviews and design write-ups.

QureAI (AI Startup)

Mar 2026 – Present

ML Infrastructure Engineer

- Integrated AWS Bedrock-hosted foundation models, including Claude, into production agentic AI workflows for managed, scalable model access as an early engineer at this AI startup.
- Improved RAG retrieval recall 61% to 89% on a 2K-query eval set by rebuilding the chunking pipeline with hybrid dense-sparse retrieval and re-ranking.

Nomura Research Institute Financial Technologies

Sep 2019 – Aug 2023

Software Engineer

- Architected and operated 25+ production Java Spring Boot microservices exposing REST APIs for 1M+ users, sustaining 99.9% uptime over 4 years on AWS (EC2, RDS, SQS).
- Drove a redesign of nightly data aggregation, cutting processing from 8 minutes to 90 seconds via indexed views, T-SQL optimization, and batch partitioning.
- Standardized failure-handling contracts (retry, timeout, propagation semantics) across 25+ services, reducing production escalations 60% in one quarter.
- Mentored engineers through code reviews and design write-ups, driving consistent engineering standards across the team.

Projects

LangFetch | tinyurl.com/LangFetch

- LangGraph agent compiling natural-language questions into dependency-ordered SQL subqueries with Pinecone vector retrieval and hybrid re-ranking; deployed as an MCP-compatible server for enterprise data queries.

RootCause AI | github.com/ganeshhgupta/rootcause-ai

- LangGraph multi-agent system diagnosing production incidents via tool-use chains; routes alerts to specialized sub-agents and synthesizes evidence-backed root cause reports, reducing mean time to resolution.

Watch-your-LLM | github.com/ganeshhgupta/watch-your-llm

- Production LLM eval platform with hallucination detection, drift tracking, rubric-based scoring, and regression alerting across model versions; FastAPI backend, PostgreSQL.

Education

University of Texas at Arlington

Aug 2023 – May 2025

MS, Computer Science, GPA 4.0/4.0

MCKV Institute of Engineering

Jul 2016 – Jul 2020

BTech, Computer Science

Skills

Languages: Python, Java, TypeScript, SQL, C++

Agentic AI & Testing: LangGraph, OpenAI Agents SDK, Prompt Engineering, LLM Evaluation, Claude Code, Codex

CI/CD & Automation: GitHub Actions, Reusable Workflows, PR Checks, Automated Testing, REST APIs

Backend & Cloud: Java, Spring Boot, Node.js, Microservices, AWS, Azure, Docker

Certifications

AWS Certified Developer Associate | GCP Professional MLOps Engineer