

GANESH GUPTA

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

Experience

Oracle

Jun 2025 – Present

Software Engineer, AI Platform

- Implement MCP server endpoints and tool-SDK integrations exposing platform capabilities to LLMs and autonomous agents, architecting multi-agent orchestration (LangGraph, OpenAI Agents SDK) at 99%+ uptime.
- Own end-to-end platform features from design doc through production validation: a deployment-gated CI pipeline enforcing guardrails and observability, blocking 15% of failed releases before shipping.
- Built full-stack developer-facing tooling (React frontend, Node.js/Python backend) and event-driven APIs on multi-tenant AWS infrastructure (Lambda, SageMaker) and Azure at sub-200ms latency.
- Designed stateful recovery via Redis checkpointing for the platform's inner development loop; interrupted workflows resume from the last committed step across distributed execution.
- Participate actively in design and code reviews, onboard new engineers, and establish CI/CD and testing standards through architecture reviews and design write-ups.

QureAI (AI Startup)

Mar 2026 – Present

ML Infrastructure Engineer

- Integrated AWS Bedrock-hosted foundation models, including Claude, into production agentic AI workflows for managed, scalable model access as an early engineer at this AI startup.
- Improved RAG retrieval recall 61% to 89% on a 2K-query eval set by rebuilding the chunking pipeline with hybrid dense-sparse retrieval and re-ranking.

Nomura Research Institute Financial Technologies

Sep 2019 – Aug 2023

Software Engineer

- Architected and operated 25+ production C++ and SQL-backed microservices in distributed systems for 1M+ users, sustaining 99.9% uptime over 4 years on AWS (EC2, RDS, SQS).
- Drove a redesign of nightly data aggregation, cutting processing from 8 minutes to 90 seconds via indexed views and batch partitioning.
- Standardized failure-handling contracts (retry, timeout, propagation semantics) across 25+ services, reducing production escalations 60% in one quarter.
- Mentored engineers through code reviews and design write-ups, driving consistent engineering standards across the team.

Projects

LangFetch | tinyurl.com/LangFetch

- LangGraph agent compiling natural-language questions into dependency-ordered SQL subqueries with Pinecone vector retrieval and hybrid re-ranking; deployed as an MCP-compatible server for enterprise data queries.

RootCause AI | github.com/ganeshhgupta/rootcause-ai

- LangGraph multi-agent system diagnosing production incidents via tool-use chains; routes alerts to specialized sub-agents and synthesizes evidence-backed root cause reports, reducing mean time to resolution.

Watch-your-LLM | github.com/ganeshhgupta/watch-your-llm

- Production LLM eval platform with hallucination detection, drift tracking, rubric-based scoring, and regression alerting across model versions; FastAPI backend, PostgreSQL.

Education

University of Texas at Arlington

Aug 2023 – May 2025

MS, Computer Science, GPA 4.0/4.0

MCKV Institute of Engineering

Jul 2016 – Jul 2020

BTech, Computer Science

Skills

Languages: Python, Java, TypeScript, SQL, C++

Agentic Dev Tooling: MCP Server Endpoints, Tool SDKs, Agentic AI, LangGraph, OpenAI Agents SDK, LLM-Integrated Tooling

Platform & Backend: React, Node.js, FastAPI, REST APIs, Microservices, CI/CD Pipelines

Cloud & Infra: AWS, GCP, Azure, Docker, Kubernetes, Terraform, Git

Certifications

AWS Certified Developer Associate | GCP Professional MLOps Engineer