

Ganesh Gupta

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

EXPERIENCE

Oracle | *Software Engineer, AI Platform*

Jun 2025 – Present

- Owned full-stack delivery of user-facing features end-to-end, from React frontend (components, hooks, state management) through Python/FastAPI backend services, at 99%+ uptime.
- Engineered a pluggable distributed executor with idempotent, retry-safe job handling across Docker and AWS Lambda, backed by an S3 artifact store eliminating redundant recomputation.
- Built a deployment-gated evaluation framework with eval-driven iteration on LLM/agent outputs, blocking 15% of failed releases before production.
- Replaced polling with SSE for live metrics streaming, cutting redundant API calls 90%, and participated in design/code reviews enforcing testability standards.

QureAI | *ML Infrastructure Engineer*

Mar 2026 – Present

- Built a model observability framework tracking drift, latency, and eval regressions in a healthcare AI production environment, enabling governance-aligned rollback workflows.
- Improved RAG retrieval recall 61% to 89% on a 2K-query eval set by rebuilding the chunking pipeline with hybrid dense-sparse retrieval and re-ranking.

Nomura Research Institute Financial Technologies | *Software Engineer*

Sep 2019 – Aug 2023

- Built and operated 25+ production microservices in Java, C++, and SQL-backed distributed systems for 1M+ users, sustaining 99.9% uptime over 3 years on AWS (EC2, RDS, SQS).
- Standardized failure-handling contracts across 25+ financial services, reducing production escalations 60% in one quarter via consistent retry, timeout, and idempotency patterns.

PROJECTS

Watch-your-LLM

github.com/ganeshhgupta/watch-your-llm

- Full-stack platform with FastAPI backend, React/TypeScript dashboard, and PostgreSQL, with eval-driven rubric-based scoring, drift detection, and CI-integrated test coverage.

LangFetch

tinyurl.com/LangFetch

- Agent decomposing multi-table questions into dependency-ordered SQL subqueries across relational and NoSQL data sources, deployed as an API server for enterprise data queries.

RootCause AI

github.com/ganeshhgupta/rootcause-ai

- Multi-agent system automating production incident diagnosis via tool-use chains and streaming LLM responses; routes alerts to specialized sub-agents.

EDUCATION

University of Texas at Arlington | *M.S. Computer Science, GPA 4.0/4.0*

Aug 2023 – May 2025

MCKV Institute of Engineering | *B.Tech Computer Science*

Jul 2016 – Jul 2020

SKILLS

Languages: Python, TypeScript, Java, C++, SQL

Full-Stack: React (hooks, state management), FastAPI, REST APIs, Async Workflows

Data: PostgreSQL, MongoDB, NoSQL, SQL/Relational Databases

AI & LLM: LLM Application Development, Streaming, Tool Calling, Eval-Driven Iteration, RAG

Cloud & DevOps: AWS, Docker, Kubernetes, CI/CD, Git

CERTIFICATIONS

AWS Certified Developer Associate | Google Cloud Professional Machine Learning Engineer