

GANESH GUPTA

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

Experience

Oracle

Jun 2025 – Present

Software Engineer, AI Platform

- Own end-to-end delivery of production data architectures combining data pipelines, ML/AI models, and user-facing interfaces, partnering directly with customer-facing stakeholders on scope and outcomes.
- Built clean, modular, well-tested Python data pipelines with automated feature engineering and Advanced SQL-backed transforms, enforcing validation and observability, blocking 15% of failed releases.
- Engineered Kafka event-stream processing across distributed services for real-time data ingestion, handling missing, duplicate, and inconsistent data before it reaches downstream consumers.
- Set CI/CD pipelines for automated testing and deployment across the platform, optimizing pipeline performance on multi-tenant AWS and Azure infrastructure.

QureAI (AI Startup)

Mar 2026 – Present

ML Infrastructure Engineer

- Built PySpark data processing jobs for large-scale dataset preparation, including cleansing, deduplication, standardization, and reconciliation, feeding downstream training and evaluation pipelines.
- Built feature engineering and data-quality pipelines (pandas, NumPy) improving retrieval recall 61% to 89% on a 2K-query eval set.

Nomura Research Institute Financial Technologies

Sep 2019 – Aug 2023

Software Engineer

- Architected and operated 25+ production C++ and SQL-backed microservices for 1M+ users at 99.9% uptime; partnered with enterprise clients to design custom integrations and translate business processes into technical workflows.
- Led data cleaning, preprocessing, and Advanced SQL storage redesign for nightly aggregation pipelines, cutting processing from 8 minutes to 90 seconds via indexed views and batch partitioning.
- Standardized failure-handling contracts (retry, timeout, propagation semantics) across 25+ services and led incident response, reducing production escalations 60% in one quarter.
- Mentored engineers through code reviews and design write-ups, driving consistent engineering standards across the team.

Projects

LangFetch | tinyurl.com/LangFetch

- LangGraph agent compiling natural-language questions into dependency-ordered SQL subqueries with Pinecone vector retrieval and hybrid re-ranking; deployed as an MCP-compatible server for enterprise data queries.

RootCause AI | github.com/ganeshhgupta/rootcause-ai

- LangGraph multi-agent system diagnosing production incidents via tool-use chains; routes alerts to specialized sub-agents and synthesizes evidence-backed root cause reports, reducing mean time to resolution.

Watch-your-LLM | github.com/ganeshhgupta/watch-your-llm

- Production LLM eval platform with hallucination detection, drift tracking, rubric-based scoring, and regression alerting across model versions; FastAPI backend, PostgreSQL.

Education

University of Texas at Arlington

Aug 2023 – May 2025

MS, Computer Science, GPA 4.0/4.0

MCKV Institute of Engineering

Jul 2016 – Jul 2020

BTech, Computer Science

Skills

Languages: Python, Advanced SQL, TypeScript, Java, C++

Data & ML Platforms: PySpark, Kafka, ETL/ELT Pipelines, ML/AI Model Integration, MLOps, Feature Engineering

Delivery & Quality: End-to-End Architecture Design, CI/CD, Automated Testing, Technical Project Delivery

Cloud & Infra: AWS (SageMaker, Lambda, EC2), Azure, GCP (Vertex AI, BigQuery), Docker, CI/CD

Certifications

AWS Certified Developer Associate | GCP Professional MLOps Engineer