

Ganesh Gupta

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

EXPERIENCE

Oracle | *Software Engineer, AI Platform*

Jun 2025 – Present

- Designed and optimized **agentic AI** workflows across **GPT**, **Claude**, **Llama**, and **Gemini**, selecting and tuning models per use case for production **quality and cost**.
- Built **evaluation harnesses** measuring LLM output quality, safety, **hallucination rates**, and grounding across agent versions; iterated on **eval results** and production usage data.
- Deployed **RAG pipeline** on **AWS SageMaker** for real-time decision support at **sub-200ms p99 latency** with **RTO under 2 minutes**.
- Built **multi-agent system** via **OpenAI Agents SDK** + **MCP** tool-use handling **millions of daily transactions** at **99%+ uptime**.

QureAI | *ML Infrastructure Engineer*

Mar 2026 – Present

- Cut **p99 inference latency 420ms to 140ms** across **50K+ daily requests** migrating synchronous serving to async **vLLM** on **GCP Cloud Run** with autoscaling.
- Built **PySpark** jobs for large-scale feature and dataset preparation feeding downstream **model training** and **evaluation** pipelines.

Nomura Research Institute Financial Technologies | *Software Engineer*

Sep 2019 – Aug 2023

- Built and operated **25+ production microservices** in **C++** and **SQL**-backed **distributed systems** for **1M+ users**, sustaining **99.9% uptime** over 4 years on **AWS** (EC2, RDS, SQS).
- Standardized failure-handling contracts across **25+ financial services**, reducing production escalations **60%** in one quarter via consistent **retry**, **timeout**, and **propagation** semantics.
- Standardized API contracts and failure handling across **25+ services**, reducing escalations **60%**; drove **root cause analysis** and **post-incident reviews** to improve long-term reliability.

PROJECTS

Watch-your-LLM

github.com/ganeshhgupta/watch-your-llm

- Production **LLM eval platform** with rubric-based scoring, **drift detection**, and latency tracking across model versions; **FastAPI**, **React/TypeScript**, **PostgreSQL**.

RootCause AI

github.com/ganeshhgupta/rootcause-ai

- LangGraph** multi-agent system diagnosing **production incidents** via tool-use chains; routes alerts to specialized sub-agents and synthesizes **root cause reports**, reducing **MTTR**.

LangFetch

tinyurl.com/LangFetch

- LangGraph** agent decomposes multi-table questions into dependency-ordered **SQL subqueries** with intermediate result passing; deployable as **MCP-compatible server**.

EDUCATION

University of Texas at Arlington | *M.S. Computer Science, GPA 4.0/4.0*

Aug 2023 – May 2025

MCKV Institute of Engineering | *B.Tech Computer Science*

Jul 2016 – Jul 2020

SKILLS

Languages: Python, SQL, TypeScript, Java, Go, HTML/CSS

ML: PyTorch, TensorFlow, Feature Engineering, Predictive Modeling, Recommendation Systems

GenAI: LLMs, Generative AI, RAG, Agentic AI, Eval Frameworks, Prompt Engineering

Data and Infra: PySpark, Airflow, AWS (SageMaker, Lambda, S3), GCP, Docker, Kubernetes

CERTIFICATIONS

AWS Certified Developer Associate | GCP Professional MLOps Engineer