

Ganesh Gupta

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

EXPERIENCE

Oracle | Senior AI Engineer, AI Platform

Jun 2025 – Present

- Designed and implemented custom agents and orchestration patterns (sequential, multi-agent, deterministic graph-based) via LangChain, LangGraph, and the OpenAI Agents SDK, with MCP servers for tool-use at 99%+ uptime.
- Engineered a cloud-native executor routing pipeline nodes to Docker containers and serverless functions via REST APIs, implementing OAuth-based auth and RBAC for secure service access.
- Designed and optimized agentic AI workflows across multiple LLMs including Gemini, GPT, Claude, and Llama, selecting and tuning models per use case for production quality and cost; use Claude Code and GitHub Copilot daily in the engineering workflow.
- Built an AI observability and model governance framework enforcing guardrails and quality/regression checks across agent and model outputs, blocking 15% of failed releases before production.

QureAI | ML Infrastructure Engineer

Mar 2026 – Present

- Built evaluation harnesses measuring LLM output quality, safety, hallucination rates, and answer grounding across agent versions, iterating rapidly based on eval results and production usage data.
- Improved RAG retrieval recall 61% to 89% on a 2K-query eval set by rebuilding the chunking pipeline with hybrid dense-sparse retrieval and re-ranking.

Nomura Research Institute Financial Technologies | Software Engineer

Sep 2019 – Aug 2023

- Architected 25+ production Java Spring Boot microservices exposing REST APIs for claims and financial transaction processing for 1M+ users at 99.9% uptime on AWS (EC2, RDS, SQS).
- Standardized REST API failure-handling contracts across 25+ Java microservices, reducing escalations 60% via consistent retry, timeout, and OAuth-aligned authentication patterns.

PROJECTS

LangFetch

tinyurl.com/LangFetch

- Full-stack app with TypeScript/Node.js backend, React dashboard, PostgreSQL, and REST API layer; integrates AI for natural language querying with OAuth-secured MCP-compatible endpoints.

Watch-your-LLM

github.com/ganeshhgupta/watch-your-llm

- Production LLM evaluation harness for quality, safety, and hallucination tracking; rubric-based scoring, regression alerting, and reusable eval components across agent versions; FastAPI, React/TypeScript, PostgreSQL.

RootCause AI

github.com/ganeshhgupta/rootcause-ai

- Multi-agent system automating production incident diagnosis via tool-use chains; routes alerts to specialized sub-agents and synthesizes evidence-backed root cause reports.

EDUCATION

University of Texas at Arlington | M.S. Computer Science, GPA 4.0/4.0

Aug 2023 – May 2025

MCKV Institute of Engineering | B.Tech Computer Science

Jul 2016 – Jul 2020

SKILLS

Languages: Java, Python, TypeScript, SQL, C++

AI/Agentic: Gemini, GPT, Claude, Claude Code, GitHub Copilot, LangChain, LangGraph, OpenAI Agents SDK, MCP, AI Observability, Model Governance

Enterprise: Spring Boot, REST APIs, Microservices, OAuth 2.0/OpenID Connect, RBAC

Cloud & DevOps: GCP (Cloud Run, Vertex AI), AWS (EC2, RDS, SQS, SageMaker), Docker, Kubernetes, CI/CD

CERTIFICATIONS

AWS Certified Developer Associate | Google Cloud Professional Machine Learning Engineer