

GANESH GUPTA

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

Experience

Oracle

Jun 2025 – Present

Software Engineer, AI Platform

- Build production GenAI/LLM applications with RAG (embeddings, semantic search, Pinecone/FAISS vector search) and prompt engineering, integrating LLMs with enterprise REST APIs.
- Own multi-step AI agent and chatbot applications (LangGraph, CrewAI, AutoGen, OpenAI Agents SDK) with tool-calling integrations at 99%+ uptime.
- Build full-stack AI application interfaces (React frontend, Python backend), partnering with product and business teams to deliver enterprise AI solutions.
- Built a deployment-gated evaluation framework enforcing guardrails and observability, blocking 15% of failed releases before production, deployed via Docker/CI-CD on AWS.

QureAI (AI Startup)

Mar 2026 – Present

ML Infrastructure Engineer

- Developed deep learning models (PyTorch, Hugging Face Transformers) with domain-specific architectures, optimizing for production inference accuracy and performance.
- Built custom evaluators for RAG/agent workflows tracking drift, latency, and regression across model versions, catching failures before they reached production.
- Improved semantic retrieval recall 61% to 89% on a 2K-query eval set by rebuilding the chunking pipeline with hybrid dense-sparse retrieval and re-ranking.

Nomura Research Institute Financial Technologies

Sep 2019 – Aug 2023

Software Engineer

- Architected and operated 25+ production C++ and SQL-backed microservices for 1M+ users, sustaining 99.9% uptime over 4 years on AWS (EC2, RDS, SQS).
- Standardized failure-handling contracts (retry, timeout, propagation semantics) across 25+ services, reducing production escalations 60% in one quarter.
- Mentored engineers through code reviews and design write-ups, driving consistent engineering standards.

Projects

LangFetch | tinyurl.com/LangFetch

- LangGraph agent with MCP-compatible tool-use compiling natural-language questions into dependency-ordered SQL subqueries with Pinecone vector retrieval and hybrid re-ranking; React dashboard, Node.js backend.

RootCause AI | github.com/ganeshhgupta/rootcause-ai

- LangGraph multi-agent system diagnosing production incidents via tool-use chains; routes alerts to specialized sub-agents and synthesizes evidence-backed root cause reports.

Watch-your-LLM | github.com/ganeshhgupta/watch-your-llm

- Production LLM eval platform with hallucination detection, drift tracking, rubric-based scoring, and regression alerting across model versions; FastAPI backend, PostgreSQL.

Education

University of Texas at Arlington

Aug 2023 – May 2025

MS, Computer Science, GPA 4.0/4.0

MCKV Institute of Engineering

Jul 2016 – Jul 2020

BTech, Computer Science

Skills

Languages: Python, TypeScript, SQL, Java, C++

GenAI & LLM: LangGraph, CrewAI, AutoGen, OpenAI Agents SDK, PyTorch, Hugging Face Transformers, Prompt Engineering

RAG & Retrieval: RAG, GraphRAG, Embeddings, Semantic Search, Pinecone, FAISS, Model/Agent Evaluation

Full-Stack & Cloud: React, REST APIs, AWS, GCP, Docker, PostgreSQL, CI/CD

Certifications

AWS Certified Developer Associate | GCP Professional MLOps Engineer