

Ganesh Gupta

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

EXPERIENCE

Oracle | *Software Engineer, AI Platform*

Jun 2025 – Present

- Built **developer-facing** agentic workflows using the **OpenAI Agents SDK** and **Model Context Protocol (MCP)** for **tool-use (function calling)** across multiple **LLMs** (GPT, Claude, Llama, Gemini).
- Designed, scaled, and maintained **large-scale infrastructure**: a pluggable distributed executor routing workloads across **Docker** and **AWS Lambda** with an **S3**-backed artifact store.
- Built a deployment-gated **evaluation** framework with automated **software quality assurance** and regression **testing**, blocking **15%** of failed releases before production.
- Used **CrewAI** and **AutoGen** alongside **LangGraph** as **orchestration frameworks** for multi-agent **GenAI** workflows handling millions of daily transactions at **99%+ uptime**.

QureAI | *ML Infrastructure Engineer*

Mar 2026 – Present

- Owned **ML infrastructure** end-to-end, model deployment through **model evaluation**, optimization, data processing, and debugging, cutting **p99 latency 420ms to 140ms** across **50K+ daily requests**.
- Improved retrieval recall **61% to 89%** on a 2K-query eval set by rebuilding the pipeline with **hybrid dense-sparse retrieval** and re-ranking.

Nomura Research Institute Financial Technologies | *Software Engineer*

Sep 2019 – Aug 2023

- Architected and operated **25+ production microservices** for **1M+ users**, sustaining **99.9% uptime** over 4 years, standardizing failure-handling contracts that cut production escalations **60%** in one quarter.

PROJECTS

RootCause AI

github.com/ganeshhgupta/rootcause-ai

- Multi-agent** system with **tool-use** chains automating production incident diagnosis; routes alerts to specialized sub-agents and synthesizes root cause reports.

LangFetch

tinyurl.com/LangFetch

- LangGraph** agent with **Pinecone** vector retrieval and hybrid re-ranking; decomposes natural-language questions into dependency-ordered SQL subqueries via **FastAPI**.

Watch-your-LLM

github.com/ganeshhgupta/watch-your-llm

- Production **LLM evaluation** platform with rubric-based scoring, hallucination detection, and **automation** for regression alerting across model versions.

EDUCATION

University of Texas at Arlington | *M.S. Computer Science, GPA 4.0/4.0*

Aug 2023 – May 2025

MCKV Institute of Engineering | *B.Tech Computer Science*

Jul 2016 – Jul 2020

SKILLS

Languages: Python, Java, C++, TypeScript, SQL

Agentic AI: OpenAI Agents SDK, MCP, Tool-Use/Function Calling, LangGraph, CrewAI, AutoGen, Multi-Modal LLMs

ML Infrastructure: Model Deployment, Model Evaluation, Optimization, Data Processing, Debugging, A/B Testing

Infra & QA: Large-Scale Infrastructure, Distributed Systems, Software QA, Testing Automation, CI/CD

Fundamentals: Data Structures & Algorithms, Operating Systems, Database Systems, Computer Networks, HTML, CSS

Cloud: AWS, GCP, Docker, Kubernetes, PostgreSQL