

# Ganesh Gupta

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

## EXPERIENCE

### Oracle | *Software Engineer, AI Platform*

Jun 2025 – Present

- Designed and optimized agentic AI workflows across multiple LLMs (GPT, Claude, Llama, Gemini), selecting and tuning models per use case for production quality and cost.
- Built a multi-agent system via OpenAI Agents SDK and MCP tool-use handling millions of daily transactions at 99%+ uptime.
- Replaced polling with SSE for live ML metrics streaming, cutting redundant API calls 90% and enabling real-time model governance dashboards.
- Engineered a pluggable distributed executor routing pipeline nodes to Docker or AWS Lambda, backed by an S3 artifact store eliminating redundant recomputation.

### QureAI | *ML Infrastructure Engineer*

Mar 2026 – Present

- Cut p99 inference latency 420ms to 140ms across 50K+ daily requests by migrating synchronous serving to async vLLM on GCP Cloud Run with autoscaling.
- Built PySpark data processing jobs for large-scale feature and dataset preparation feeding downstream model training and evaluation pipelines.

### University of Texas at Arlington | *Graduate Teaching Assistant*

Aug 2024 – May 2025

- Taught and mentored graduate students in a machine learning course through office hours and lab sessions, translating model architectures, training methodology, and evaluation techniques into applied understanding.
- Developed multi-modal AI models using PyTorch and TensorFlow with attention mechanisms and domain-specific neural network architectures; built evaluation frameworks measuring model quality and performance.

### Nomura Research Institute Financial Technologies | *Software Engineer*

Sep 2019 – Aug 2023

- Built and operated 25+ production microservices in Java, C++, and SQL-backed distributed systems for 1M+ users, sustaining 99.9% uptime over 3 years on AWS (EC2, RDS, SQS).
- Cut nightly financial data aggregation blocking downstream settlement from 8 minutes to 90 seconds by redesigning SQL with indexed views and batch partitioning.
- Standardized failure-handling contracts across 25+ financial services, reducing production escalations 60% in one quarter via consistent retry, timeout, and propagation semantics.

## PROJECTS

### Watch-your-LLM

[github.com/ganeshhgupta/watch-your-llm](https://github.com/ganeshhgupta/watch-your-llm)

- Production LLM eval platform with rubric-based scoring, drift detection, and latency tracking across model versions; FastAPI backend, React/TypeScript dashboard, PostgreSQL.

### RootCause AI

[github.com/ganeshhgupta/rootcause-ai](https://github.com/ganeshhgupta/rootcause-ai)

- LangGraph multi-agent system diagnosing production incidents via tool-use chains; routes alerts to specialized sub-agents and synthesizes root cause reports, reducing mean time to resolution.

### LangFetch

[tinyurl.com/LangFetch](https://tinyurl.com/LangFetch)

- LangGraph agent decomposing multi-table questions into dependency-ordered SQL subqueries with intermediate result passing; deployable as an MCP-compatible server for enterprise data queries.

## EDUCATION

University of Texas at Arlington | *M.S. Computer Science, GPA 4.0/4.0*

Aug 2023 – May 2025

MCKV Institute of Engineering | *B.Tech Computer Science*

Jul 2016 – Jul 2020

## SKILLS

**Languages:** Python, Java, C++, TypeScript, SQL

**AI/ML:** LLMs (GPT, Claude, Llama, Gemini), PyTorch, TensorFlow, Agentic AI, Multi-Agent Systems (LangGraph, OpenAI Agents SDK, MCP), Model Evaluation & Observability

**Data:** PySpark, PostgreSQL, SQL/Relational Databases, NoSQL

**Cloud & DevOps:** AWS (EC2, S3, Lambda), GCP (Cloud Run), Docker, Kubernetes, CI/CD

## CERTIFICATIONS

AWS Certified Developer Associate | Google Cloud Professional Machine Learning Engineer