

Ganesh Gupta

817-822-1131 | ganeshg12ta@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

EXPERIENCE

Oracle | *Software Engineer, AI Platform*

Jun 2025 – Present

- Architected and shipped **multi-agent GenAI** systems using **LangGraph**, **CrewAI**, and **AutoGen** as orchestration frameworks, taking them from prototype to production at **millions of daily transactions** and **99%+ uptime**.
- Deployed a **RAG** pipeline on **AWS SageMaker** for real-time decision support at **sub-200ms p99** latency with RTO under 2 minutes.
- Designed **agent state management** with **Redis checkpointing**, so interrupted multi-step workflows resume from the last committed step across distributed execution.
- Tuned agentic workflows across **GPT**, **Claude**, **Llama**, and **Gemini** on **LLM-native metrics** (quality, latency, **cost per request**), with **MCP** tool-use and **SSE** streaming of live agent output to users.

QureAI | *ML Infrastructure Engineer*

Mar 2026 – Present

- Improved **RAG** retrieval recall **61% to 89%** on a 2K-query eval set by rebuilding chunking with **hybrid dense-sparse retrieval** over **vector databases** and re-ranking.
- Cut p99 inference latency **420ms to 140ms** across **50K+ daily requests** by re-architecting serving to async **vLLM** on **GCP Cloud Run** with autoscaling.

Nomura Research Institute Financial Technologies | *Software Engineer*

Sep 2019 – Aug 2023

- Partnered with **enterprise clients** to design custom integrations, translating business processes into technical workflows across **25+ production microservices** serving **1M+ users** at **99.9% uptime**.
- Standardized failure-handling contracts across 25+ services and drove root-cause analysis with stakeholders, reducing production escalations **60%** in one quarter.

PROJECTS

LangFetch

tinyurl.com/LangFetch

- LangGraph** agent with **Pinecone** vector retrieval and hybrid re-ranking; decomposes natural-language questions into dependency-ordered SQL subqueries via **FastAPI**, deployable as an **MCP** server.

RootCause AI

github.com/ganeshhgupta/rootcause-ai

- Multi-agent** system with **ReAct**-style tool-use chains for incident diagnosis; routes alerts to specialized sub-agents and synthesizes root-cause reports.

Watch-your-LLM

github.com/ganeshhgupta/watch-your-llm

- LLM evaluation** and **tracing** platform with rubric-based scoring, hallucination detection, and regression alerting across model versions.

EDUCATION

University of Texas at Arlington | *M.S. Computer Science, GPA 4.0/4.0*

Aug 2023 – May 2025

MCKV Institute of Engineering | *B.Tech Computer Science*

Jul 2016 – Jul 2020

SKILLS

Languages: Python, Java, TypeScript, SQL, C++

Agentic AI: LangGraph, CrewAI, AutoGen, OpenAI Agents SDK, MCP, ReAct, Tool Use, Multi-Agent Orchestration

RAG & Data: Vector Databases (Pinecone), Hybrid Retrieval, Re-ranking, GraphRAG (Neo4j), Structured/Unstructured Pipelines, PySpark

Cloud & LLMOps: GCP (Vertex AI, Cloud Run, BigQuery), AWS (SageMaker, Bedrock, Lambda), Docker, Kubernetes, vLLM, Evaluation, Tracing

Fundamentals: Data Structures & Algorithms, Operating Systems, Database Systems, Computer Networks

Certifications: AWS Certified Developer Associate, Google Cloud Professional Machine Learning Engineer