

Ganesh Gupta

817-822-1131 | iamgsl0rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

EXPERIENCE

Oracle | *Software Engineer, AI Platform*

Jun 2025 – Present

- Developed infrastructure services in **Go**, **Python**, and **C++** for **GPU** scheduling, resource allocation, and fault-tolerant workload placement across **AWS**, **GCP**, and **CoreWeave compute** fleets operating under capacity constraints.
- Designed **GPU** fleet orchestration for training and inference workloads, evaluating scheduling, utilization, reliability, and capacity tradeoffs through cross-team **architecture** reviews.
- Debugged production failures spanning compute nodes, **network** dependencies, schedulers, and serving services, driving fixes through implementation, testing, code review, and production rollout.
- Built **Prometheus/Grafana** observability and autoscaling for inference serving, correlating compute, **network**, and **storage** telemetry to identify fleet-wide performance regressions.

QureAI | *ML Infrastructure Engineer*

Mar 2026 – Present

- Built and operated **Kubernetes**-based **GPU** infrastructure across multiple cloud providers, developing **Go** and **Python** tooling for workload scheduling and fault-tolerant placement during node failures and provider disruptions.
- Reduced **p99 serving latency** from **420ms to 140ms** across 50K+ daily requests through asynchronous execution redesign, profiling, and systems-level performance optimization.

Nomura Research Institute Financial Technologies | *Software Engineer*

Sep 2019 – Aug 2023

- Designed and operated **25+ distributed C++/Java** services supporting **1M+ users** on **AWS EC2**, **RDS**, and **SQS**, sustaining **99.9% availability** across four years of production ownership.
- Led **architecture** decisions for shared platform components and standardized service contracts, retry/timeout behavior, and failure-handling patterns across 25+ services, reducing production escalations **60%**.
- Diagnosed production incidents across application, database, **network**, and downstream-service boundaries using logs, metrics, and root-cause analysis; codified recurring failure fixes into shared platform standards and code-review checks.
- Redesigned a nightly **SQL** aggregation pipeline blocking downstream settlement, reducing runtime from **8 minutes to 90 seconds** through indexed views and batch partitioning while eliminating row-level lock contention.

PROJECTS

Watch-your-LLM

github.com/ganeshhgupta/watch-your-llm

- Built a full-stack model **evaluation** platform with automated regression suites integrated into **CI/CD**, using **FastAPI** and **PostgreSQL** to track quality changes across deployments.

RootCause AI

github.com/ganeshhgupta/rootcause-ai

- Built a multi-agent incident-diagnosis system that orchestrates specialized tools and services to collect evidence, analyze production failures, and synthesize root-cause reports.

EDUCATION

University of Texas at Arlington | *M.S. Computer Science, GPA 4.0/4.0*

Aug 2023 – May 2025

MCKV Institute of Engineering | *B.Tech Computer Science*

Jul 2016 – Jul 2020

SKILLS

Languages: Go, Python, C++, Java, SQL

Distributed Systems: Distributed Services, Scheduling, Fault Tolerance, Resource Allocation, Concurrency, Service Reliability

Infrastructure: Linux, GPU Compute Infrastructure, Networking, Storage Systems, Observability, Performance Profiling

Cloud & DevOps: AWS, GCP, CoreWeave, Kubernetes, Docker, Terraform, CI/CD, Prometheus, Grafana

Data Systems: PostgreSQL, AWS RDS, SQS, SQL Performance Optimization

Certifications: AWS Certified Developer Associate, Google Cloud Professional Machine Learning Engineer