

Ganesh Gupta

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

EXPERIENCE

Oracle | *Software Engineer, AI Platform*

Jun 2025 – Present

- Own multi-step agent workflows end-to-end using LangGraph, CrewAI, AutoGen, and the OpenAI Agents SDK, with tool-calling integrations and stateful agent memory via Redis checkpointing at 99%+ uptime.
- Design ML infrastructure spanning model deployment, evaluation, and data processing: combined GraphRAG with vector-based RAG for grounded retrieval on complex enterprise queries.
- Built a deployment-gated evaluation framework enforcing guardrails, model quality checks, and observability across agent/model outputs, blocking 15% of failed releases before production.
- Architected scalable platform APIs (Python, TypeScript) enabling product teams to build, test, and launch GenAI-powered applications end-to-end.

QureAI | *ML Infrastructure Engineer*

Mar 2026 – Present

- Developed multi-modal computer vision models (PyTorch) for medical imaging analysis, running hands-on model training and fine-tuning on multi-GPU clusters with compression and quantization for efficient inference.
- Built custom evaluators for RAG/agent workflows tracking drift, latency, and regression across model versions, catching failures before they reached production.
- Improved semantic retrieval recall 61% to 89% on a 2K-query eval set by rebuilding the chunking pipeline with hybrid dense-sparse retrieval and re-ranking.

Nomura Research Institute Financial Technologies | *Software Engineer*

Sep 2019 – Aug 2023

- Architected and operated 25+ production C++ and SQL-backed microservices for 1M+ users, sustaining 99.9% uptime over 4 years on AWS (EC2, RDS, SQS).
- Drove a redesign of nightly data aggregation, cutting processing from 8 minutes to 90 seconds via indexed views and batch partitioning.
- Standardized failure-handling contracts (retry, timeout, propagation semantics) across 25+ services, reducing production escalations 60% in one quarter.

PROJECTS

LangFetch

tinyurl.com/LangFetch

- LangGraph agent with MCP-compatible tool-use compiling natural-language questions into dependency-ordered SQL subqueries with Pinecone vector retrieval and hybrid re-ranking; deployed as an MCP server for enterprise data queries.

Watch-your-LLM

github.com/ganeshhgupta/watch-your-llm

- Production LLM eval platform with hallucination detection, drift tracking, rubric-based scoring, and regression alerting across model versions; FastAPI backend, PostgreSQL.

RootCause AI

github.com/ganeshhgupta/rootcause-ai

- LangGraph multi-agent system diagnosing production incidents via tool-use chains; routes alerts to specialized sub-agents and synthesizes evidence-backed root cause reports, reducing mean time to resolution.

EDUCATION

University of Texas at Arlington | *M.S. Computer Science, GPA 4.0/4.0*

Aug 2023 – May 2025

MCKV Institute of Engineering | *B.Tech Computer Science*

Jul 2016 – Jul 2020

SKILLS

Languages: Python, TypeScript, SQL, Java, C++

GenAI & ML Infra: PyTorch, LangGraph, CrewAI, AutoGen, OpenAI Agents SDK, Model Training/Fine-Tuning (Multi-GPU), Model Compression/Quantization

Retrieval, Vision & Evaluation: RAG, GraphRAG, Computer Vision, Multi-Modal Learning, Model/Agent Evaluation, Guardrails, Observability

Cloud & Platform: AWS, GCP, Docker, Kubernetes, PostgreSQL, Redis, CI/CD

CERTIFICATIONS

AWS Certified Developer Associate | Google Cloud Professional Machine Learning Engineer