

GANESH GUPTA

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

Experience

Oracle

Jun 2025 – Present

Software Engineer, AI Platform

- Lead design of agentic AI systems (LangGraph, CrewAI, AutoGen, OpenAI Agents SDK) combining GraphRAG with vector-based retrieval for grounded reasoning over enterprise data.
- Own model quality: built a deployment-gated evaluation framework measuring accuracy, latency, and regression, enforcing guardrails before production release.
- Partner directly with product and business stakeholders to translate ambiguous problems into deployed AI/ML solutions (Python, React/TypeScript).
- Set engineering standards for CI/CD and containerized deployment across the AI platform, working in Agile sprints.

QureAI (AI Startup)

Mar 2026 – Present

ML Infrastructure Engineer

- Ran hands-on model training and fine-tuning (PyTorch) on multi-GPU clusters with domain-specific architectures, applying compression and quantization for efficient inference.
- Built custom evaluators measuring drift, latency, and regression across model versions, catching quality degradation before it reached production.
- Improved semantic retrieval recall 61% to 89% on a 2K-query eval set by rebuilding the chunking pipeline with hybrid dense-sparse retrieval and re-ranking.

Nomura Research Institute Financial Technologies

Sep 2019 – Aug 2023

Software Engineer

- Architected and operated 25+ production C++ and SQL-backed microservices for 1M+ users, sustaining 99.9% uptime over 4 years on AWS (EC2, RDS, SQS).
- Drove a redesign of nightly data aggregation, cutting processing from 8 minutes to 90 seconds via indexed views and batch partitioning.
- Standardized failure-handling contracts (retry, timeout, propagation semantics) across 25+ services, reducing production escalations 60% in one quarter.

Projects

Watch-your-LLM | github.com/ganeshhgupta/watch-your-llm

- Production LLM eval platform with hallucination detection, drift tracking, rubric-based scoring, and regression alerting across model versions; FastAPI backend, PostgreSQL.

RootCause AI | github.com/ganeshhgupta/rootcause-ai

- LangGraph multi-agent system diagnosing production incidents via tool-use chains; routes alerts to specialized sub-agents and synthesizes evidence-backed root cause reports, reducing mean time to resolution.

LangFetch | tinyurl.com/LangFetch

- LangGraph agent with MCP-compatible tool-use compiling natural-language questions into dependency-ordered SQL subqueries with Pinecone vector retrieval and hybrid re-ranking; deployed as an MCP server for enterprise data queries.

Education

University of Texas at Arlington

Aug 2023 – May 2025

MS, Computer Science, GPA 4.0/4.0

MCKV Institute of Engineering

Jul 2016 – Jul 2020

BTech, Computer Science

Skills

Languages: Python, TypeScript, SQL, Java, C++

ML & Data Science: PyTorch, Model Training/Fine-Tuning (Multi-GPU), Model Compression/Quantization, Model Evaluation

Agentic AI & Retrieval: LangGraph, CrewAI, AutoGen, OpenAI Agents SDK, RAG, GraphRAG, Guardrails

Cloud & Platform: AWS, GCP, Docker, Kubernetes, PostgreSQL, CI/CD

Certifications

AWS Certified Developer Associate | GCP Professional MLOps Engineer