

GANESH GUPTA

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

Experience

Oracle

Jun 2025 – Present

Software Engineer, AI Platform

- Own technical design for the agentic AI serving platform, architecting multi-agent orchestration (LangGraph, LangChain, OpenAI Agents SDK, MCP) with usability, reliability, and scalability as first-class design goals at 99%+ uptime.
- Built a deployment-gated model evaluation pipeline enforcing guardrails, governance, and observability across the platform, leading build-vs-buy tooling decisions and blocking 15% of failed releases before production.
- Designed stateful recovery via Redis checkpointing for the serving pipeline on multi-tenant AWS (Lambda, SageMaker) and Azure; interrupted workflows resume from the last committed step, cutting incident recovery time.
- Partner with ML engineers, data engineers, and product; set standards for CI/CD, testing, and responsible-AI guardrails, mentoring engineers through architecture reviews.

QureAI (AI Startup)

Mar 2026 – Present

ML Infrastructure Engineer

- Built a model observability framework tracking drift, latency, and evaluation regressions in production, enabling governance-aligned rollback workflows as an early engineer at this AI startup.
- Improved RAG retrieval recall 61% to 89% on a 2K-query eval set by rebuilding the chunking pipeline with hybrid dense-sparse retrieval and re-ranking.

Nomura Research Institute Financial Technologies

Sep 2019 – Aug 2023

Software Engineer

- Architected and operated 25+ production C++ and SQL-backed microservices in distributed systems for 1M+ users, sustaining 99.9% uptime over 4 years on AWS (EC2, RDS, SQS).
- Led data cleaning, preprocessing, and storage redesign for nightly aggregation pipelines, cutting processing from 8 minutes to 90 seconds via indexed views and batch partitioning.
- Standardized failure-handling contracts (retry, timeout, propagation semantics) across 25+ services and led incident response, reducing production escalations 60% in one quarter.
- Mentored engineers through code reviews and design write-ups, driving consistent engineering standards across the team.

Projects

LangFetch | tinyurl.com/LangFetch

- LangGraph agent compiling natural-language questions into dependency-ordered SQL subqueries with Pinecone vector retrieval and hybrid re-ranking; deployed as an MCP-compatible server for enterprise data queries.

RootCause AI | github.com/ganeshhgupta/rootcause-ai

- LangGraph multi-agent system diagnosing production incidents via tool-use chains; routes alerts to specialized sub-agents and synthesizes evidence-backed root cause reports, reducing mean time to resolution.

Watch-your-LLM | github.com/ganeshhgupta/watch-your-llm

- Production LLM eval platform with hallucination detection, drift tracking, rubric-based scoring, and regression alerting across model versions; FastAPI backend, PostgreSQL.

Education

University of Texas at Arlington

Aug 2023 – May 2025

MS, Computer Science, GPA 4.0/4.0

MCKV Institute of Engineering

Jul 2016 – Jul 2020

BTech, Computer Science

Skills

Languages: Python, Java, TypeScript, SQL, C++

ML Platform & MLOps: LangGraph, LangChain, MCP, RAG, Model Observability, Evaluation Pipelines, AI Governance, Data Pipelines

Backend & Security: FastAPI, REST APIs, Microservices, PostgreSQL, RBAC, Guardrails

Cloud & Infra: AWS, GCP (Vertex AI, BigQuery), Azure, Docker, Kubernetes, CI/CD

Certifications

AWS Certified Developer Associate | GCP Professional MLOps Engineer