

Ganesh Gupta

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

EXPERIENCE

Oracle | *Software Engineer, AI Platform*

Jun 2025 – Present

- Owned end-to-end production path for AI-backed services: deployment, serving, release automation, monitoring, and rollback, integrating LLM capabilities into backend product workflows via clean API interfaces.
- Engineered a pluggable distributed executor routing asynchronous jobs across Docker and AWS Lambda with an S3 artifact store, applying idempotency and failure-handling to eliminate redundant recomputation.
- Deployed a RAG pipeline on AWS SageMaker for real-time decision support at sub-200ms p99 latency with RTO under 2 minutes, with graceful degradation on model-backed component failure.
- Built a multi-agent system via OpenAI Agents SDK and MCP tool-use handling millions of daily transactions at 99%+ uptime, with structured logging and dashboards for operability.
- Replaced polling with SSE for live metrics streaming, cutting redundant API calls 90%; participated in design and code reviews enforcing testability and production-readiness standards.

QureAI | *ML Infrastructure Engineer*

Mar 2026 – Present

- Built a model observability framework tracking drift, latency, and eval regressions in production, enabling versioning and governance-aligned rollback workflows for deployed models.
- Cut p99 inference latency 420ms to 140ms across 50K+ daily requests by migrating synchronous model serving to async vLLM on GCP Cloud Run with autoscaling.

Nomura Research Institute Financial Technologies | *Software Engineer*

Sep 2019 – Aug 2023

- Built and operated 25+ production microservices in Java, C++, and SQL-backed distributed systems for 1M+ users in a regulated financial environment, sustaining 99.9% uptime over 3 years on AWS (EC2, RDS, SQS).
- Standardized failure-handling contracts across 25+ financial services, reducing production escalations 60% in one quarter via consistent retry, timeout, and propagation semantics.
- Cut nightly financial data aggregation blocking downstream settlement from 8 minutes to 90 seconds by redesigning SQL with indexed views and batch partitioning.

PROJECTS

Watch-your-LLM

github.com/ganeshhgupta/watch-your-llm

- Production LLM eval and observability platform with rubric-based scoring, drift detection, and model-version latency tracking; FastAPI backend, React/TypeScript dashboard, PostgreSQL.

RootCause AI

github.com/ganeshhgupta/rootcause-ai

- LangGraph multi-agent system automating production incident diagnosis via tool-use chains; routes alerts to specialized sub-agents and synthesizes root cause reports, reducing manual toil.

Distributed KV Store / VPC Control Plane Simulator

github.com/ganeshhgupta/distributed-kv-store

- Raft-consensus distributed KV store handling concurrent writes with linearizable reads and log compaction; companion VPC control plane simulator modeling subnet routing and security group enforcement.

EDUCATION

University of Texas at Arlington | *M.S. Computer Science, GPA 4.0/4.0*

Aug 2023 – May 2025

MCKV Institute of Engineering | *B.Tech Computer Science*

Jul 2016 – Jul 2020

SKILLS

Languages: Python, TypeScript, Java, C++, SQL

Distributed Systems: APIs, Async Processing (Lambda, SQS), Retries/Idempotency/Graceful Degradation, Observability, Structured Logging

AI/ML Integration: LLMs (OpenAI, Anthropic), RAG, Agentic AI, Model Evaluation & Versioning, OpenAI Agents SDK, MCP

Cloud & DevOps: AWS (EC2, S3, SageMaker, Lambda, SQS, RDS), GCP (Cloud Run), Docker, Kubernetes, CI/CD, PostgreSQL

CERTIFICATIONS

AWS Certified Developer Associate | Google Cloud Professional Machine Learning Engineer