

GANESH GUPTA

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

Experience

Oracle

Jun 2025 – Present

Software Engineer, AI Platform

- Technical anchor for the agentic AI platform, architecting multi-agent orchestration and tool-use reasoning using LangGraph, LangChain, and OpenAI Agents SDK with MCP tool integration at 99%+ uptime.
- Built a deployment-gated evaluation pipeline enforcing guardrails, AI governance, and observability around agent outputs, blocking 15% of failed releases.
- Built production-grade Python/SQL data pipelines and curated datasets (Snowflake) powering RAG, analytics, and automation workflows exposed as REST APIs at sub-200ms latency.
- Partner directly with product and business stakeholders to turn real problems into deployed AI solutions, translating between technical and non-technical audiences.
- Set engineering standards for CI/CD, testing, and secure development practices across the platform, working in Agile sprints with Git-based version control.

QureAI (AI Startup)

Mar 2026 – Present

ML Infrastructure Engineer

- Built a model observability and governance framework tracking drift, latency, and evaluation regressions in production, enabling monitored, auditable rollback workflows.
- Improved semantic search/RAG retrieval recall 61% to 89% on a 2K-query eval set by rebuilding the chunking pipeline with hybrid dense-sparse retrieval and re-ranking.

Nomura Research Institute Financial Technologies

Sep 2019 – Aug 2023

Software Engineer

- Architected and operated 25+ production C++ and SQL-backed microservices for 1M+ users, sustaining 99.9% uptime over 4 years on AWS (EC2, RDS, SQS).
- Drove a redesign of nightly data aggregation, cutting processing from 8 minutes to 90 seconds via indexed views and batch partitioning.
- Standardized failure-handling contracts (retry, timeout, propagation semantics) across 25+ services, reducing production escalations 60% in one quarter.
- Mentored engineers through code reviews and design write-ups, driving consistent engineering standards.

Projects

LangFetch | tinyurl.com/LangFetch

- LangGraph agent compiling natural-language questions into dependency-ordered SQL subqueries with Pinecone vector retrieval and hybrid re-ranking; deployed as an MCP-compatible server for enterprise data queries.

RootCause AI | github.com/ganeshhgupta/rootcause-ai

- LangGraph multi-agent system diagnosing production incidents via tool-use chains; routes alerts to specialized sub-agents and synthesizes evidence-backed root cause reports, reducing mean time to resolution.

Watch-your-LLM | github.com/ganeshhgupta/watch-your-llm

- Production LLM eval platform with hallucination detection, drift tracking, rubric-based scoring, and regression alerting across model versions; FastAPI backend, PostgreSQL.

Education

University of Texas at Arlington

Aug 2023 – May 2025

MS, Computer Science, GPA 4.0/4.0

MCKV Institute of Engineering

Jul 2016 – Jul 2020

BTech, Computer Science

Skills

Languages: Python, SQL, Java, TypeScript, C++

AI/ML & RAG: Prompt Engineering, RAG, LLM APIs, Agentic AI, LangGraph, LangChain, Model Governance & Observability

Data & Backend: Snowflake, Data Pipelines, REST APIs, Vector DBs (Pinecone, FAISS), Guardrails

Cloud & DevOps: Azure, AWS, GCP, Docker, Git, CI/CD

Certifications

AWS Certified Developer Associate | GCP Professional MLOps Engineer