

Ganesh Gupta

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

EXPERIENCE

Oracle | *Software Engineer, AI Platform*

Jun 2025 – Present

- Designed and optimized agentic AI workflows in Python across multiple LLMs (GPT, Claude, Llama, Gemini), with tool/function calling and structured outputs via the OpenAI Agents SDK and MCP.
- Built a deployment-gated evaluation framework enforcing guardrails, governance, and quality/regression checks across model outputs, blocking 15% of failed releases before production.
- Deployed a RAG pipeline on AWS SageMaker for real-time decision support at sub-200ms p99 latency with RTO under 2 minutes, troubleshooting and optimizing performance in production.
- Owned backend services and APIs (Python/FastAPI) handling millions of daily transactions at 99%+ uptime.

QureAI | *ML Infrastructure Engineer*

Mar 2026 – Present

- Improved RAG retrieval recall 61% to 89% on a 2K-query eval set by rebuilding the chunking and document ingestion pipeline with hybrid dense-sparse retrieval, embeddings, and re-ranking.
- Built evaluation harnesses measuring LLM output quality, safety, and hallucination rates across agent versions, with observability and regression testing integrated into the release process.

Nomura Research Institute Financial Technologies | *Software Engineer*

Sep 2019 – Aug 2023

- Built and operated 25+ production microservices in Java, C++, and SQL-backed distributed systems for 1M+ users across financial transaction and institutional lending workflows, sustaining 99.9% uptime over 3 years on AWS.
- Standardized failure-handling contracts across 25+ financial services, reducing production escalations 60% in one quarter via consistent retry, timeout, and propagation semantics.

PROJECTS

Watch-your-LLM

github.com/ganeshhgupta/watch-your-llm

- Production LLM eval and observability platform with rubric-based scoring, drift detection, hallucination tracking, and regression alerting across model versions; FastAPI backend, PostgreSQL.

LangFetch

tinyurl.com/LangFetch

- LangGraph agent with MCP-compatible tool-use compiling natural-language questions into dependency-ordered SQL subqueries with Pinecone vector retrieval and hybrid re-ranking.

RootCause AI

github.com/ganeshhgupta/rootcause-ai

- Multi-agent system automating production incident diagnosis via tool-use chains; routes alerts to specialized sub-agents and synthesizes root cause reports.

EDUCATION

University of Texas at Arlington | *M.S. Computer Science, GPA 4.0/4.0*

Aug 2023 – May 2025

MCKV Institute of Engineering | *B.Tech Computer Science*

Jul 2016 – Jul 2020

SKILLS

Languages: Python, Java, TypeScript, C++, SQL

Enterprise GenAI: LLMs (GPT, Claude, Llama, Gemini), RAG, Embeddings, Vector Databases, Semantic/Hybrid Search, Re-Ranking

Agentic AI: Tool/Function Calling, Structured Outputs, LangChain, LangGraph, OpenAI Agents SDK, MCP

LLMOps & Governance: Evaluation, Observability, Regression Testing, Guardrails, Governance

Cloud & DevOps: AWS (EC2, S3, SageMaker), GCP, Docker, Kubernetes, CI/CD, PostgreSQL

CERTIFICATIONS

AWS Certified Developer Associate | Google Cloud Professional Machine Learning Engineer