

Ganesh Gupta

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

EXPERIENCE

Oracle | *Software Engineer, AI Platform*

Jun 2025 – Present

- Architected a production **multi-agent system** on the **OpenAI Agents SDK** with **MCP tool-use**, automating enterprise workflows across **millions of daily transactions** at **99%+ uptime**.
- Led **agentic AI** orchestration across **LangGraph**, **CrewAI**, and **AutoGen**, matching multi-agent patterns to each use case.
- Owned a **RAG pipeline** on **AWS SageMaker** for real-time decision support at **sub-200ms p99 latency** and **RTO under 2 minutes**; added **GraphRAG** to ground answers on multi-hop queries.
- Tuned **GenAI** workflows across **GPT**, **Claude**, **Llama**, and **Gemini** via **prompt engineering**, picking models by quality and cost.
- Drove stateful agent recovery via **Redis checkpointing**; interrupted **agent workflows** resume from the last committed step.

QureAI | *ML Infrastructure Engineer*

Mar 2026 – Present

- Raised **RAG retrieval recall** from **61% to 89%** on a **2K-query eval set** by rebuilding chunking with **hybrid dense-sparse retrieval** and **re-ranking**.
- Built healthcare **agentic AI** on **AWS Bedrock Agents**, **Knowledge Bases**, and **Guardrails** for managed retrieval and safety.

Nomura Research Institute Financial Technologies | *Software Engineer*

Sep 2019 – Aug 2023

- Built and operated **25+ production microservices** in **Python**, **C++**, and **SQL-backed distributed systems** for **1M+ users** at **99.9% uptime** on **AWS** (EC2, RDS, SQS).
- Integrated **ML models** (**scikit-learn**, **TensorFlow**) into production workflows with data pipelines processing **50K+ records**.
- Standardized **REST API** and failure-handling contracts across **25+ services**, cutting production escalations **60%** in one quarter; mentored engineers on reliability practices.
- Cut nightly settlement-blocking aggregation from **8 minutes to 90 seconds** with **SQL indexed views** and **batch partitioning**.

PROJECTS

RootCause AI

github.com/ganeshhgupta/rootcause-ai

- LangGraph multi-agent system** diagnosing incidents via **tool-use chains** and specialized sub-agents, reducing **MTTR**.

LangFetch

tinyurl.com/LangFetch

- LangGraph agent** decomposes multi-table questions into dependency-ordered **SQL subqueries**; deployable as an **MCP server**.

Watch-your-LLM

github.com/ganeshhgupta/watch-your-llm

- LLM evaluation** platform with rubric scoring, **drift detection**, and latency tracking; **FastAPI**, **React**, **PostgreSQL**.

EDUCATION

University of Texas at Arlington | *M.S. Computer Science, GPA 4.0/4.0*

Aug 2023 – May 2025

MCKV Institute of Engineering | *B.Tech Computer Science*

Jul 2016 – Jul 2020

SKILLS

Languages: Python, TypeScript, JavaScript, Java, SQL, C++, HTML/CSS, Data Structures and Algorithms

GenAI & Agentic AI: LangGraph, LangChain, OpenAI Agents SDK, CrewAI, AutoGen, MCP, RAG, Prompt Engineering

LLMs & Data: GPT, Claude, Llama, Gemini, AWS Bedrock, Pinecone, FAISS, Hugging Face, vLLM, Redis, PostgreSQL

Cloud & MLOps: AWS (SageMaker, Lambda, S3, EC2), GCP (Vertex AI, Cloud Run), Docker, Kubernetes, FastAPI, CI/CD

CERTIFICATIONS

AWS Certified Developer Associate | GCP Professional MLOps Engineer