

GANESH GUPTA

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

Experience

Oracle

Jun 2025 – Present

Software Engineer, AI Platform

- Provision and operate multi-cloud GPU clusters (AWS, GCP, CoreWeave) with NVIDIA/CUDA hardware and InfiniBand/RoCE networking, validating GPU node health across heterogeneous hardware for training and inference workloads.
- Write infrastructure/systems-level code in Go, Rust, and C++ for Kubernetes-based scheduling, resource allocation, and fault-tolerant placement logic under hardware-scarcity constraints.
- Built observability and autoscaling for vLLM inference serving via Prometheus/Grafana, and run distributed training workloads via Slurm-based HPC scheduling, maintaining fleet utilization through node loss.
- Own CI/CD, Terraform-based infrastructure as code, and containerized deployment (Docker/Kubernetes), writing runbooks and architecture documentation for the platform.

QureAI (AI Startup)

Mar 2026 – Present

ML Infrastructure Engineer

- Ran hands-on model training/fine-tuning on multi-GPU clusters with compression and quantization for efficient production inference.
- Built PySpark data processing jobs and custom evaluators for ML/agent workflows, tracking drift, latency, and regression across model versions before production.

Nomura Research Institute Financial Technologies

Sep 2019 – Aug 2023

Software Engineer

- Architected and operated 25+ production C++ and SQL-backed microservices for 1M+ users in a regulated financial services environment, sustaining 99.9% uptime over 4 years on AWS (EC2, RDS, SQS).
- Standardized failure-handling contracts (retry, timeout, propagation semantics) across 25+ distributed services, reducing production escalations 60% in one quarter.
- Mentored engineers through code reviews and design write-ups, driving consistent engineering standards.

Projects

Distributed KV Store & VPC Control Plane Simulator | github.com/ganeshhgupta/distributed-kv-store

- Built a distributed key-value store in Go/C++ implementing Raft consensus, linearizable reads, and log compaction; simulated VPC control-plane semantics for network partition and leader-election testing.

Watch-your-LLM | github.com/ganeshhgupta/watch-your-llm

- Production ML/LLM eval platform with drift tracking, rubric-based scoring, and regression alerting across model versions; FastAPI backend, PostgreSQL.

RootCause AI | github.com/ganeshhgupta/rootcause-ai

- Multi-agent system diagnosing production infrastructure incidents via tool-use chains; routes alerts to specialized sub-agents and synthesizes evidence-backed root cause reports.

Education

University of Texas at Arlington

Aug 2023 – May 2025

MS, Computer Science, GPA 4.0/4.0

MCKV Institute of Engineering

Jul 2016 – Jul 2020

BTech, Computer Science

Skills

Languages: Python, Go, C++, Rust, SQL

ML Infra & GPU: GPU Clusters (NVIDIA/CUDA), Slurm HPC Scheduling, Model Training/Fine-Tuning (Multi-GPU), PySpark

Platform & DevOps: Kubernetes, Docker, Terraform, CI/CD, Prometheus/Grafana, Distributed Systems

Cloud: AWS (EC2, S3, RDS), GCP, CoreWeave

Certifications

AWS Certified Developer Associate | GCP Professional MLOps Engineer