

GANESH GUPTA

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

Experience

Oracle

Jun 2025 – Present

Software Engineer, AI Platform

- Owned end-to-end design of a multi-agent system using OpenAI Agents SDK and MCP tool-calling; prompt design, reasoning chains, and tool integration with internal systems, powering RAG-based decision support for millions of daily transactions at 99%+ uptime.
- Hardened the system from working to reliable: built a deployment-gated evaluation pipeline enforcing automated compliance, quality, latency, and drift checks, blocking 15% of releases that failed validation before reaching production.
- Architected AWS serverless infrastructure (Lambda, S3, DynamoDB, API Gateway, SageMaker) for real-time agentic inference at sub-200ms latency, supporting high-throughput production traffic at scale.

QureAI (AI Startup)

Jan 2024 – May 2025

ML Infrastructure Engineer

- Diagnosed a user-facing RAG retrieval gap directly from real usage signals, then owned the fix end-to-end: redesigned chunking and prompt strategy, lifting recall 61% to 89% and cutting the underlying complaint pattern from 22% to 6% of sessions.
- Led migration of synchronous model serving to async vLLM with autoscaling as an early engineer at this AI startup, cutting p99 inference latency 420ms to 140ms across 50K+ daily requests.

Nomura Research Institute Financial Technologies

Aug 2020 – Aug 2023

Software Engineer

- Architected and operated 25+ production Java microservices processing regulated financial transactions for 1M+ users, sustaining 99.9% uptime over 3 years on AWS (EC2, RDS, SQS).
- Drove a redesign of nightly financial data aggregation, cutting processing from 8 minutes to 90 seconds via indexed views and batch partitioning, in a domain requiring auditable, traceable transaction records.
- Led standardization of failure-handling contracts across 25+ services in partnership with cross-functional teams through code reviews and design write-ups, reducing production escalations 60% in one quarter.

Projects

Watch-your-LLM | github.com/ganeshhgupta/watch-your-llm

- Production LLM eval platform with hallucination detection, drift tracking, rubric-based scoring, and regression alerting across model versions; FastAPI backend, PostgreSQL, built to keep agentic systems trustworthy at volume.

RootCause AI | github.com/ganeshhgupta/rootcause-ai

- LangGraph multi-agent system diagnosing production incidents via tool-use chains; routes alerts to specialized sub-agents and synthesizes evidence-backed, traceable root cause reports.

Education

University of Texas at Arlington

Aug 2023 – May 2025

MS, Computer Science, GPA 4.0/4.0

MCKV Institute of Engineering

Jul 2016 – Jul 2020

BTech, Computer Science

Skills

Languages: Python, Java, TypeScript, SQL

ML/AI Engineering: RAG, Agentic AI, LangGraph, LangChain, OpenAI Agents SDK, MCP, Vector Databases (Pinecone, FAISS), Prompt Engineering, Evals, Hallucination/Drift Detection, vLLM

Backend & APIs: FastAPI, REST APIs, Node.js, PostgreSQL

MLOps & Cloud: AWS (Lambda, EC2, S3, RDS, SQS, SageMaker), GCP (Cloud Run), Docker, Kubernetes, CI/CD, Production ML Deployment

Certifications

AWS Certified Developer Associate | GCP Professional MLOps Engineer