

Ganesh Gupta

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

EXPERIENCE

Oracle | *Software Engineer, AI Platform*

Jun 2025 – Present

- Provisioned and operated **multi-cloud GPU clusters** (AWS, GCP, CoreWeave) with **NVIDIA/CUDA** hardware and **InfiniBand/RoCE** networking for long-running training and **high-availability inference**.
- Wrote systems-level code in **Go, Rust**, and **C++** for **scheduling, resource allocation**, and **fault-tolerant placement** under GPU-scarcity constraints.
- Built observability and **autoscaling** for **vLLM** serving via **Prometheus/Grafana**; ran distributed training via **Slurm** HPC scheduling, keeping utilization up through node loss.
- Built **multi-agent system** via **OpenAI Agents SDK + MCP** tool-use handling **millions of daily transactions** at **99%+ uptime**.

QureAI | *ML Infrastructure Engineer*

Mar 2026 – Present

- Cut **p99 inference latency 420ms to 140ms** across **50K+ daily requests** migrating synchronous serving to async **vLLM** on **GCP Cloud Run** with autoscaling.
- Built **model observability** framework tracking **drift**, latency, and **eval regressions** in production, enabling governance-aligned **rollback workflows**.

Nomura Research Institute Financial Technologies | *Software Engineer*

Sep 2019 – Aug 2023

- Built and operated **25+ production microservices** in **C++** and SQL-backed **distributed systems** for **1M+ users**, sustaining **99.9% uptime** over 4 years on **AWS** (EC2, RDS, SQS).
- Cut nightly financial data aggregation blocking downstream **settlement** from **8 minutes to 90 seconds** by redesigning SQL with **indexed views** and **batch partitioning**.
- Standardized failure-handling contracts across **25+ financial services**, reducing production escalations **60%** in one quarter via consistent **retry, timeout, and propagation** semantics.

PROJECTS

Watch-your-LLM

github.com/ganeshhgupta/watch-your-llm

- Production **LLM eval platform** with rubric-based scoring, **drift detection**, and latency tracking across model versions; **FastAPI**, **React/TypeScript**, **PostgreSQL**.

RootCause AI

github.com/ganeshhgupta/rootcause-ai

- LangGraph** multi-agent system diagnosing **production incidents** via tool-use chains; routes alerts to specialized sub-agents and synthesizes **root cause reports**, reducing **MTTR**.

LangFetch

tinyurl.com/LangFetch

- LangGraph** agent decomposes multi-table questions into dependency-ordered **SQL subqueries** with intermediate result passing; deployable as **MCP-compatible server**.

EDUCATION

University of Texas at Arlington | *M.S. Computer Science, GPA 4.0/4.0*

Aug 2023 – May 2025

MCKV Institute of Engineering | *B.Tech Computer Science*

Jul 2016 – Jul 2020

SKILLS

Languages: Java, Python, Go, SQL, TypeScript, HTML/CSS

ML and Recommenders: PyTorch, Recommender Systems, Ranking and Retrieval, A/B Testing, LLM Fine-tuning

GPU and Serving: GPU Fleets, vLLM, Model Serving, Prometheus/Grafana, Autoscaling

Data and Infra: PySpark, AWS, GCP, Kubernetes, Docker, PostgreSQL, Redis

CERTIFICATIONS

AWS Certified Developer Associate | GCP Professional MLOps Engineer