

Ganesh Gupta

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

EXPERIENCE

Oracle | *Software Engineer, AI Platform*

Jun 2025 – Present

- Designed and optimized agentic AI workflows in Python across multiple LLMs including Anthropic Claude, GPT, Llama, and Gemini, selecting and tuning models per use case for production quality and cost.
- Built a multi-agent system via the OpenAI Agents SDK and MCP (Model Context Protocol) servers for tool-use, handling millions of daily transactions at 99%+ uptime.
- Deployed a RAG pipeline on AWS SageMaker for real-time decision support at sub-200ms p99 latency with RTO under 2 minutes, working directly with downstream consumers to scope requirements.
- Owned full-stack delivery of backend services and APIs (Python/FastAPI, React/TypeScript), replacing polling with SSE for live metrics streaming and cutting redundant API calls 90%.

QureAI | *ML Infrastructure Engineer*

Mar 2026 – Present

- Built a model observability framework tracking drift, latency, and eval regressions in production, enabling governance-aligned rollback workflows.
- Improved RAG retrieval recall 61% to 89% on a 2K-query eval set by rebuilding the chunking pipeline with hybrid dense-sparse retrieval and re-ranking.

Nomura Research Institute Financial Technologies | *Software Engineer*

Sep 2019 – Aug 2023

- Built and operated 25+ production microservices in Java, C++, and SQL-backed distributed systems for 1M+ users, sustaining 99.9% uptime over 3 years on AWS (EC2, RDS, SQS).
- Standardized failure-handling contracts across 25+ financial services, reducing production escalations 60% in one quarter via consistent retry, timeout, and propagation semantics.

PROJECTS

LangFetch

tinyurl.com/LangFetch

- LangGraph agent decomposing multi-table questions into dependency-ordered SQL subqueries with intermediate result passing; deployable as an MCP-compatible server for client-facing data queries.

RootCause AI

github.com/ganeshhgupta/rootcause-ai

- LangGraph multi-agent system automating production incident diagnosis via tool-use chains; routes alerts to specialized sub-agents and synthesizes root cause reports, reducing manual toil.

Watch-your-LLM

github.com/ganeshhgupta/watch-your-llm

- Production LLM eval and observability platform with rubric-based scoring, drift detection, and latency tracking across model versions; FastAPI backend, React/TypeScript dashboard, PostgreSQL.

EDUCATION

University of Texas at Arlington | *M.S. Computer Science, GPA 4.0/4.0*

Aug 2023 – May 2025

MCKV Institute of Engineering | *B.Tech Computer Science*

Jul 2016 – Jul 2020

SKILLS

Languages: Python, TypeScript, Java, C++, SQL

AI/ML: LLMs (Anthropic Claude, GPT, Llama, Gemini), RAG, LangChain, LangGraph, Agentic AI, OpenAI Agents SDK, MCP Servers

Full-Stack: FastAPI, React, REST APIs, PostgreSQL, MongoDB

Cloud & DevOps: AWS (EC2, S3, SageMaker), GCP (Cloud Run), Docker, Kubernetes, CI/CD

CERTIFICATIONS

AWS Certified Developer Associate | Google Cloud Professional Machine Learning Engineer