

GANESH GUPTA

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

Experience

Oracle

Jun 2025 – Present

Software Engineer, AI Platform

- Architect enterprise AI solutions on Azure: designed a multi-agent system (OpenAI Agents SDK, LangGraph, CrewAI, AutoGen, MCP tool-calling) with planning, reasoning, and memory management for production workflow automation.
- Built RAG pipelines with vector search and embedding-based retrieval (Pinecone, FAISS), plus GraphRAG improving LLM response grounding over standard vector RAG for complex, multi-hop queries.
- Designed multi-step LLM evaluation tasks and an experimentation pipeline analyzing model performance and regressions across versions, identifying where and why model behavior degrades before release.
- Hardened the system from working to reliable: built an AI evaluation, monitoring, and observability pipeline enforcing automated compliance, quality, and drift checks, blocking 15% of releases that failed validation.

QureAI (Healthcare AI Startup)

Mar 2026 – Present

ML Infrastructure Engineer

- Diagnosed a model-quality gap directly from real usage signals and ran controlled experiments to fix it end-to-end: redesigned retrieval, lifting recall 61% to 89% and cutting the complaint pattern from 22% to 6% of sessions.
- Led migration of PyTorch-backed model serving to async vLLM with autoscaling, cutting p99 inference latency 420ms to 140ms across 50K+ daily requests.

Nomura Research Institute Financial Technologies

Sep 2019 – Aug 2023

Software Engineer

- Architected and operated 25+ production Java microservices processing regulated financial transactions for 1M+ users, sustaining 99.9% uptime over 4 years on AWS (EC2, RDS, SQS).
- Drove a redesign of nightly financial data aggregation, cutting processing from 8 minutes to 90 seconds via indexed views and batch partitioning, in a domain requiring auditable, traceable transaction records.
- Led standardization of failure-handling contracts across 25+ services in partnership with cross-functional teams through code reviews and design write-ups, reducing production escalations 60% in one quarter.

Projects

Watch-your-LLM | github.com/ganeshhgupta/watch-your-llm

- Production LLM eval platform with hallucination detection, drift tracking, rubric-based scoring, and regression alerting across model versions; FastAPI backend, PostgreSQL, built to keep agentic systems trustworthy at volume.

RootCause AI | github.com/ganeshhgupta/rootcause-ai

- LangGraph multi-agent system diagnosing production incidents via tool-use chains; routes alerts to specialized sub-agents and synthesizes evidence-backed, traceable root cause reports.

LangFetch | tinyurl.com/LangFetch

- LangGraph agent with Pinecone vector retrieval and hybrid re-ranking decomposes multi-table questions into dependency-ordered SQL subqueries; deployable as MCP-compatible server.

Education

University of Texas at Arlington

Aug 2023 – May 2025

MS, Computer Science, GPA 4.0/4.0

MCKV Institute of Engineering

Jul 2016 – Jul 2020

BTech, Computer Science

Skills

Languages: Python, SQL, Java, TypeScript

ML/AI Engineering: PyTorch, Reinforcement Learning, Model Evaluation & Experimentation, RAG, Agentic AI, LangGraph, OpenAI Agents SDK, MCP, Vector Databases (Pinecone, FAISS)

Backend & APIs: FastAPI, REST APIs, PostgreSQL

Cloud & MLOps: Azure, AWS, GCP, Docker, Kubernetes, CI/CD

Certifications

AWS Certified Developer Associate | GCP Professional MLOps Engineer