

Ganesh Gupta

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

EXPERIENCE

Oracle | *Software Engineer, AI Platform*

Jun 2025 – Present

- Designed and optimized agentic AI workflows across multiple LLMs (GPT, Claude, Llama, Gemini), selecting and tuning models per use case for production quality and cost.
- Replaced polling with SSE for live ML metrics streaming, cutting redundant API calls 90% and enabling real-time model governance dashboards.
- Built a deployment-gated evaluation and performance-tooling platform used across the engineering team, with rubric-based benchmarking that blocked 15% of regressions before release.
- Drove design and code reviews enforcing style and testability standards, mentoring engineers and triaging production issues affecting service reliability.
- Built a multi-agent system via OpenAI Agents SDK and MCP tool-use handling millions of daily transactions at 99%+ uptime.

QureAI | *ML Infrastructure Engineer*

Mar 2026 – Present

- Improved RAG retrieval recall 61% to 89% on a 2K-query eval set by rebuilding the chunking pipeline with hybrid dense-sparse retrieval and re-ranking.
- Cut p99 inference latency 420ms to 140ms across 50K+ daily requests by migrating synchronous serving to async vLLM on GCP Cloud Run with autoscaling.

Nomura Research Institute Financial Technologies | *Software Engineer*

Sep 2019 – Aug 2023

- Architected and operated 25+ production microservices for 1M+ users, sustaining 99.9% uptime over 3 years on AWS (EC2, RDS, SQS).
- Standardized failure-handling contracts across 25+ financial services, reducing production escalations 60% in one quarter via consistent retry, timeout, and propagation semantics, and mentored engineers on the new standard.
- Cut nightly financial data aggregation blocking downstream settlement from 8 minutes to 90 seconds by redesigning SQL with indexed views and batch partitioning.

PROJECTS

Watch-your-LLM

github.com/ganeshhgupta/watch-your-llm

- Production eval platform with rubric-based scoring, drift detection, and latency tracking across model versions; FastAPI backend, React/TypeScript dashboard, PostgreSQL.

RootCause AI

github.com/ganeshhgupta/rootcause-ai

- LangGraph multi-agent system automating production incident diagnosis via tool-use chains; routes alerts to specialized sub-agents and synthesizes root cause reports, reducing manual toil.

LangFetch

tinyurl.com/LangFetch

- LangGraph agent decomposing multi-table questions into dependency-ordered SQL subqueries with intermediate result passing; deployable as an MCP-compatible server for enterprise data queries.

EDUCATION

University of Texas at Arlington | *M.S. Computer Science, GPA 4.0/4.0*

Aug 2023 – May 2025

MCKV Institute of Engineering | *B.Tech Computer Science*

Jul 2016 – Jul 2020

SKILLS

Languages: Python, TypeScript, JavaScript, Java, C++, SQL

AI/ML: LLMs (GPT, Claude, Llama, Gemini), RAG, LangChain, LangGraph, Agentic AI, OpenAI Agents SDK, MCP

Frontend & Backend: React, FastAPI, Node.js, REST APIs

Cloud & Infra: AWS (EC2, RDS, SQS, Lambda, S3), GCP (Cloud Run), Docker, Kubernetes, PostgreSQL

CERTIFICATIONS

AWS Certified Developer Associate | Google Cloud Professional Machine Learning Engineer