

GANESH GUPTA

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

Experience

Oracle

Jun 2025 – Present

Software Engineer, AI Platform

- Set technical direction for agentic AI initiatives end-to-end: architected a multi-agent LLM system (OpenAI Agents SDK, MCP tool-calling) integrating production LLMs into user-facing decision flows at 99%+ uptime.
- Built data quality and observability tooling: automated evaluation, drift checks, and dataset validation gating model releases, blocking 15% of releases that failed validation before reaching production.
- Mentored engineers and drove code-review and design standards across the team, establishing ownership of the platform end-to-end from architecture through production operation.

QureAI

Mar 2026 – Present

ML Infrastructure Engineer

- Took an ML system from inception to business impact: led migration of production model serving to async vLLM with autoscaling, cutting p99 inference latency 420ms to 140ms across 50K+ daily requests at scale.
- Diagnosed a model-quality gap directly from real usage signals and owned the fix end-to-end, lifting retrieval recall 61% to 89% and cutting the complaint pattern from 22% to 6% of sessions.

Nomura Research Institute Financial Technologies

Sep 2019 – Aug 2023

Software Engineer

- Architected and operated 25+ production C++ and Java microservices in a distributed, fault-tolerant system for 1M+ users, sustaining 99.9% uptime over 4 years, debugging and optimizing code for quality and performance.
- Led complex technical efforts and set technical direction across teams: standardized failure-handling contracts across 25+ services, reducing production escalations 60% in one quarter through architecture reviews.

Projects

Distributed KV Store / VPC Control Plane Simulator | github.com/ganeshhgupta/distributed-kv-store

- Built a Raft-consensus key-value store in-house with leader election, log replication, and log compaction; companion VPC control plane simulator modeling subnet routing and security group enforcement for fault-tolerant distributed coordination.

Watch-your-LLM | github.com/ganeshhgupta/watch-your-llm

- Production performance and eval tooling with rubric-based scoring, drift detection, and per-call latency benchmarking across model versions; FastAPI backend, PostgreSQL.

RootCause AI | github.com/ganeshhgupta/rootcause-ai

- LangGraph multi-agent system diagnosing production incidents via tool-use chains; routes alerts to specialized sub-agents and synthesizes evidence-backed root cause reports.

Education

University of Texas at Arlington

Aug 2023 – May 2025

MS, Computer Science, GPA 4.0/4.0

MCKV Institute of Engineering

Jul 2016 – Jul 2020

BTech, Computer Science

Skills

Languages: C++, Java, Python, JavaScript, SQL

ML & AI at Scale: PyTorch, vLLM, LLM Integration & Evaluation, RAG, Agentic AI, LangGraph, OpenAI Agents SDK, MCP, Prompt Engineering

Distributed Systems: Fault Tolerance, High Availability, Autoscaling, Data Pipelines, PySpark, Prometheus, Grafana

Cloud & Infra: AWS (EC2, Lambda, S3, DynamoDB), GCP (Cloud Run), Docker, Kubernetes, Git, CI/CD

Certifications

AWS Certified Developer Associate | GCP Professional MLOps Engineer