

Ganesh Gupta

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

EXPERIENCE

Oracle | *Software Engineer, AI Platform*

Jun 2025 – Present

- Provisioned and operated multi-cloud GPU clusters (AWS, GCP, CoreWeave) with NVIDIA/CUDA hardware and InfiniBand/RoCE networking, supporting training and inference workloads on shared fleet capacity.
- Wrote infrastructure and systems-level code in Go, Rust, and C++ for scheduling, resource allocation, and fault-tolerant placement logic under real GPU-scarcity constraints.
- Built observability and autoscaling for vLLM inference serving using Prometheus and Grafana, and ran distributed training workloads via Slurm-based HPC scheduling, keeping the fleet utilized through node loss and provider hiccups.
- Built a deployment-gated evaluation and diagnostic-tooling platform used across the engineering team, with rubric-based benchmarking that blocked 15% of regressions before release.

QureAI | *ML Infrastructure Engineer*

Mar 2026 – Present

- Built a model observability framework tracking drift, latency, and eval regressions in production, enabling governance-aligned rollback workflows.
- Cut p99 inference latency 420ms to 140ms across 50K+ daily requests by migrating synchronous serving to async vLLM on GCP Cloud Run with autoscaling.

Nomura Research Institute Financial Technologies | *Software Engineer*

Sep 2019 – Aug 2023

- Built and operated 25+ production microservices in Java, C++, and SQL-backed distributed systems for 1M+ users, sustaining 99.9% uptime over 3 years on AWS (EC2, RDS, SQS).
- Standardized failure-handling contracts across 25+ financial services, reducing production escalations 60% in one quarter via consistent retry, timeout, and propagation semantics.

PROJECTS

RootCause AI

github.com/ganeshhgupta/rootcause-ai

- LangGraph multi-agent system driving AI-native, LLM-guided diagnostics to reduce time-to-root-cause on production incidents; routes alerts to specialized sub-agents via tool-use chains.

Watch-your-LLM

github.com/ganeshhgupta/watch-your-llm

- Production eval/diagnostic platform with rubric-based scoring, drift detection, and latency tracking across model versions and serving infrastructure; FastAPI backend, PostgreSQL.

Distributed KV Store / VPC Control Plane Simulator

github.com/ganeshhgupta/distributed-kv-store

- Raft-consensus distributed KV store handling concurrent writes with linearizable reads and log compaction; companion VPC control plane simulator modeling subnet routing and security group enforcement.

EDUCATION

University of Texas at Arlington | *M.S. Computer Science, GPA 4.0/4.0*

Aug 2023 – May 2025

MCKV Institute of Engineering | *B.Tech Computer Science*

Jul 2016 – Jul 2020

SKILLS

Languages: C++, Go, Rust, Python, Java, SQL

GPU & HPC: NVIDIA/CUDA, InfiniBand/RoCE, Multi-Cloud GPU Orchestration (AWS, GCP, CoreWeave), Slurm

Tooling & Diagnostics: Eval/Regression-Gating Harnesses, Model Observability, Drift/Latency Monitoring, LLM-Guided Root-Cause Diagnostics

Infra: Prometheus, Grafana, Docker, Kubernetes, Distributed Systems, Fault Tolerance

CERTIFICATIONS

AWS Certified Developer Associate | Google Cloud Professional Machine Learning Engineer