

GANESH GUPTA

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

Experience

Oracle

Jun 2025 – Present

Software Engineer, AI Platform

- Set technical direction for a multi-agent GenAI system using LangGraph and the OpenAI Agents SDK with MCP-based agent orchestration, leading design through deployment and mentoring engineers.
- Architected a RAG retrieval system on AWS SageMaker applying IR techniques (indexing, dense-sparse ranking) for real-time decision support at sub-200ms p99 latency.
- Integrated Claude Code and Codex into team workflows, redesigning engineering processes to drive measurable efficiency gains.
- Designed stateful agent recovery via Redis checkpointing; interrupted workflows resume from the last committed step across distributed execution.
- Drove a deployment-gated evaluation pipeline implementing responsible AI practices (bias mitigation, quality review) to block 15% of failed releases.

QureAI (AI Startup)

Mar 2026 – Present

ML Infrastructure Engineer

- Designed and implemented a Neo4j knowledge graph modeling domain entities and relationships, applying graph theory to power graph-aware retrieval and reasoning in production.
- Improved RAG retrieval recall 61% to 89% on a 2K-query eval set by rebuilding the chunking pipeline with hybrid dense-sparse retrieval, ranking, and re-ranking.

Nomura Research Institute Financial Technologies

Sep 2019 – Aug 2023

Software Engineer

- Architected and operated 25+ production C++ and SQL-backed microservices in distributed systems for 1M+ users, sustaining 99.9% uptime over 4 years on AWS (EC2, RDS, SQS).
- Drove a redesign of nightly data aggregation, cutting processing from 8 minutes to 90 seconds via indexed views and batch partitioning.
- Standardized failure-handling contracts (retry, timeout, propagation semantics) across 25+ services, reducing production escalations 60% in one quarter.
- Mentored engineers through code reviews and design write-ups, driving consistent engineering standards across the team.

Projects

LangFetch | tinyurl.com/LangFetch

- LangGraph agent compiling natural-language questions into dependency-ordered SQL subqueries with Pinecone vector retrieval and hybrid re-ranking; deployed as an MCP-compatible server for enterprise data queries.

RootCause AI | github.com/ganeshhgupta/rootcause-ai

- LangGraph multi-agent system diagnosing production incidents via tool-use chains; routes alerts to specialized sub-agents and synthesizes evidence-backed root cause reports, reducing mean time to resolution.

Watch-your-LLM | github.com/ganeshhgupta/watch-your-llm

- Production LLM eval platform with hallucination detection, drift tracking, rubric-based scoring, and regression alerting across model versions; FastAPI backend, PostgreSQL.

Education

University of Texas at Arlington

Aug 2023 – May 2025

MS, Computer Science, GPA 4.0/4.0

MCKV Institute of Engineering

Jul 2016 – Jul 2020

BTech, Computer Science

Skills

Languages: Python, C++, Java, TypeScript, SQL

ML & IR: PyTorch, TensorFlow, scikit-learn, RAG, Vector Search & Ranking, Knowledge Graphs

Agentic AI & GenAI: Agentic AI, LangGraph, LangChain, OpenAI Agents SDK, Prompt Engineering, MCP

Cloud & MLOps: AWS (SageMaker, Lambda), Azure, GCP, Docker, Kubernetes, CI/CD

Certifications

AWS Certified Developer Associate | GCP Professional MLOps Engineer