

Ganesh Gupta

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

EXPERIENCE

Oracle | *Software Engineer, AI Platform*

Jun 2025 – Present

- Built production **AI** systems: multi-agent **LLM** workflows (**LangGraph**, **OpenAI Agents SDK**, **MCP**) at **millions of daily transactions** and **99%+ uptime**.
- Deployed a **RAG** pipeline on **AWS SageMaker** for real-time decision support at **sub-200ms p99** latency with RTO under 2 minutes.
- Built **MCP** tool-use and **function calling** across **GPT**, **Claude**, **Llama**, and **Gemini**, tuning model choice per task on quality, latency, and **cost per request**.
- Shipped a deployment-gated **evaluation** framework with automated regression **testing**, blocking **15%** of failed releases before production.

QureAI | *ML Infrastructure Engineer*

Mar 2026 – Present

- Owned **ML infrastructure** end to end, from model deployment to evaluation and debugging, cutting **p99 latency 420ms to 140ms** across **50K+ daily requests**.
- Improved retrieval recall **61% to 89%** on a 2K-query eval set by rebuilding chunking with **hybrid dense-sparse retrieval** and re-ranking.

Nomura Research Institute Financial Technologies | *Software Engineer*

Sep 2019 – Aug 2023

- Built and operated **25+ production microservices** for **1M+ users** on **AWS** (EC2, RDS, SQS), sustaining **99.9% uptime** over four years.
- **Profiled compute and memory bottlenecks** across **Python/C++** services and optimized batching and **concurrency**, raising throughput **2.4x** and cutting memory **27%**.
- Designed **Oracle PL/SQL** and **SQL** data layers behind financial services, tuning queries and schemas for low-latency reads at production scale.

PROJECTS

Watch-your-LLM

github.com/ganeshhgupta/watch-your-llm

- **LLM evaluation** platform with rubric-based scoring, drift detection, and latency tracking across model versions; **FastAPI**, **PostgreSQL**, regression alerting.

LangFetch

tinyurl.com/LangFetch

- **LangGraph** agent decomposing natural-language questions into dependency-ordered SQL subqueries with hybrid retrieval and re-ranking; deployable as an **MCP** server.

RootCause AI

github.com/ganeshhgupta/rootcause-ai

- **Multi-agent** incident-diagnosis system with **tool-use** chains that collects evidence from logs and metrics and synthesizes root-cause reports.

EDUCATION

University of Texas at Arlington | *M.S. Computer Science, GPA 4.0/4.0*

Aug 2023 – May 2025

MCKV Institute of Engineering | *B.Tech Computer Science*

Jul 2016 – Jul 2020

SKILLS

Languages: Python, C++, Java, Go, SQL

AI/ML: LLMs, RAG, Retrieval & Ranking, Evaluation, PyTorch, vLLM

Systems: Online Serving, Low-Latency Services, Distributed Systems, Kafka

Cloud & DevOps: AWS, GCP, Kubernetes, Docker, Terraform, CI/CD, Prometheus, Grafana

Fundamentals: Data Structures & Algorithms, Operating Systems, Database Systems, Computer Networks, HTML, CSS

Certifications: AWS Certified Developer Associate, Google Cloud Professional Machine Learning Engineer