

Ganesh Gupta

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

EXPERIENCE

Oracle | *Software Engineer, AI Platform*

Jun 2025 – Present

- Architected **agentic decision-support** systems with **multi-step reasoning, planning**, and **multi-agent** decomposition (**LangGraph, CrewAI, AutoGen**), running at **millions of daily transactions** and **99%+ uptime**.
- Turned "make the agent better" into a measurable problem: built a deployment-gated **agent evaluation** framework with representative test sets and regression **benchmarks**, blocking **15%** of failed releases before production.
- Designed **stateful workflows** with **Redis checkpointing** and **structured outputs**, so interrupted multi-step agent runs resume from the last committed step with consistent state.
- Built adaptive **model routing** across **GPT, Claude, Llama**, and **Gemini** with **MCP** tool use, choosing models per task on measured quality, latency, and **cost per request**.

QureAI | *ML Infrastructure Engineer*

Mar 2026 – Present

- Raised retrieval recall **61% to 89%** on a **2K-query evaluation set** by rebuilding chunking with **hybrid dense-sparse retrieval** and re-ranking, validating gains against a fixed held-out benchmark.
- Owned **ML infrastructure** end to end, from deployment to **evaluation** and debugging, cutting **p99 latency 420ms to 140ms** across **50K+ daily requests**.

Nomura Research Institute Financial Technologies | *Software Engineer*

Sep 2019 – Aug 2023

- Built **data pipelines** and **APIs** behind **financial services** decision workflows across **25+ production microservices** for **1M+ users** at **99.9% uptime**.
- Designed **Oracle PL/SQL** and **SQL** data layers and partnered with **business stakeholders** to formalize ambiguous requirements into measurable system behavior.

PROJECTS

Watch-your-LLM

github.com/ganeshhgupta/watch-your-llm

- **Agent/LLM evaluation** platform with rubric-based scoring, **drift detection**, and regression alerting; compares model versions on performance distributions, not anecdotes.

RootCause AI

github.com/ganeshhgupta/rootcause-ai

- **Multi-agent** decision system with **ReAct** tool-use chains and automated critique that collects evidence, ranks failure hypotheses, and synthesizes root-cause reports.

LangFetch

tinyurl.com/LangFetch

- **LangGraph** planner decomposing natural-language questions into dependency-ordered SQL subqueries with structured outputs; deployable as an **MCP** server.

EDUCATION

University of Texas at Arlington | *M.S. Computer Science, GPA 4.0/4.0*

Aug 2023 – May 2025

MCKV Institute of Engineering | *B.Tech Computer Science*

Jul 2016 – Jul 2020

SKILLS

Languages: Python, SQL, Java, TypeScript, C++

Agentic AI: Multi-Step Reasoning & Planning, Multi-Agent Systems, Agent Evaluation & Benchmarking, Memory & Context, Model Routing, LangGraph, CrewAI, AutoGen, MCP

ML & Statistics: Statistical Modelling, Experiment Design, Evaluation Datasets, PyTorch, Reinforcement Learning, RAG, Hybrid Retrieval

Engineering: AI-Native Development (Codex, Claude Code coding agents), APIs, Data Pipelines, FastAPI, PostgreSQL, Redis, Docker, Kubernetes, AWS, GCP

Fundamentals: Data Structures & Algorithms, Operating Systems, Database Systems, Computer Networks, HTML, CSS

Certifications: AWS Certified Developer Associate, Google Cloud Professional Machine Learning Engineer