

Ganesh Gupta

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

EXPERIENCE

Oracle | *Software Engineer, AI Platform*

Jun 2025 – Present

- Built AI agents following Plan-Act-Observe-Improve patterns via LangGraph and the OpenAI Agents SDK, with MCP-based execution harnesses for tool, API, and enterprise system integration.
- Designed context pipelines combining RAG, memory, and prompt strategies across multiple LLMs (GPT, Claude, Llama, Gemini), tuning models per use case for production quality and cost.
- Owned backend services and APIs (Python/FastAPI) handling millions of daily transactions at 99%+ uptime, with monitoring and performance optimization built in from the start.
- Built a deployment-gated evaluation framework enforcing guardrails and quality/regression checks across agent outputs, blocking 15% of failed releases before production.

QureAI | *ML Infrastructure Engineer*

Mar 2026 – Present

- Built a model observability framework tracking drift, latency, and eval regressions in production, with structured logging and monitoring across agent and model versions.
- Improved RAG retrieval recall 61% to 89% on a 2K-query eval set by rebuilding the chunking pipeline with hybrid dense-sparse retrieval over a vector store and re-ranking.

Nomura Research Institute Financial Technologies | *Software Engineer*

Sep 2019 – Aug 2023

- Built and operated 25+ production microservices in Java, C++, and SQL-backed distributed systems for 1M+ users, sustaining 99.9% uptime over 3 years on AWS (EC2, RDS, SQS).
- Standardized failure-handling contracts across 25+ financial services, reducing production escalations 60% in one quarter via consistent retry, timeout, and propagation semantics.

PROJECTS

RootCause AI

github.com/ganeshhgupta/rootcause-ai

- LangGraph multi-agent system automating production incident diagnosis via tool-use chains; routes alerts to specialized sub-agents and synthesizes root cause reports.

LangFetch

tinyurl.com/LangFetch

- LangGraph agent with MCP-compatible tool-use decomposing multi-table questions into dependency-ordered SQL subqueries; deployed as an API server for enterprise data queries.

Watch-your-LLM

github.com/ganeshhgupta/watch-your-llm

- Production LLM eval and observability platform with rubric-based scoring, drift detection, and latency tracking across model versions; FastAPI backend, React/TypeScript dashboard, PostgreSQL.

EDUCATION

University of Texas at Arlington | *M.S. Computer Science, GPA 4.0/4.0*

Aug 2023 – May 2025

MCKV Institute of Engineering | *B.Tech Computer Science*

Jul 2016 – Jul 2020

SKILLS

Languages: Python, TypeScript, Java, C++, SQL

Agentic AI: LLMs (GPT, Claude, Llama, Gemini), LangChain, LangGraph, AI Agents, RAG, Vector Databases, Context Pipelines

Backend: FastAPI, REST APIs, Enterprise System Integration, Monitoring & Logging

Cloud & DevOps: AWS (EC2, S3, SageMaker), GCP (Cloud Run), Docker, Kubernetes, CI/CD, PostgreSQL

CERTIFICATIONS

AWS Certified Developer Associate | Google Cloud Professional Machine Learning Engineer