

GANESH GUPTA

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

Experience

Oracle

Jun 2025 – Present

Software Engineer, AI Platform

- Design and deploy an agentic platform in Python using LangGraph, the OpenAI Agents SDK, CrewAI, and AutoGen with MCP tool-calling, architecting multi-agent orchestration for tool use, memory, and decision-making at 99%+ uptime.
- Deployed and operated the platform on Kubernetes across AWS and Azure, integrating agentic services with client-facing APIs and data platforms while partnering directly with product stakeholders.
- Drove cross-team technical strategy for a deployment-gated evaluation pipeline governing agent and AI-generated output quality, raising the engineering bar by blocking 15% of failed releases before production.

QureAI (AI Startup)

Mar 2026 – Present

ML Infrastructure Engineer

- Designed and deployed a RAG pipeline integrating enterprise data sources, improving retrieval recall 61% to 89% on a 2K-query eval set through hybrid dense-sparse retrieval and re-ranking over structured and unstructured data.
- Architected a model observability platform tracking drift, latency, and evaluation regressions across production agent/model outputs, establishing governance-aligned rollback workflows as an early engineer at this AI startup.

Nomura Research Institute Financial Technologies

Sep 2019 – Aug 2023

Software Engineer

- Led the architecture and design of 25+ production C++ and SQL-backed microservices in distributed systems for 1M+ users, sustaining 99.9% uptime over 4 years on AWS (EC2, RDS, SQS).
- Drove a redesign of nightly data aggregation and query logic, cutting processing from 8 minutes to 90 seconds via indexed views and batch partitioning across the data layer.
- Led standardization of failure-handling contracts (retry, timeout, propagation semantics) across 25+ services, mentoring engineers through code reviews and design write-ups, reducing production escalations 60% in one quarter.

Projects

LangFetch | tinyurl.com/LangFetch

- LangGraph agent compiling natural-language questions into dependency-ordered SQL subqueries with Pinecone vector retrieval and hybrid re-ranking; deployed as an MCP-compatible server for enterprise data queries.

RootCause AI | github.com/ganeshhgupta/rootcause-ai

- LangGraph multi-agent system diagnosing production incidents via tool-use chains; routes alerts to specialized sub-agents and synthesizes evidence-backed root cause reports, reducing mean time to resolution.

Education

University of Texas at Arlington

Aug 2023 – May 2025

MS, Computer Science, GPA 4.0/4.0

MCKV Institute of Engineering

Jul 2016 – Jul 2020

BTech, Computer Science

Skills

Languages: Python, Go, SQL, TypeScript, Java

Agentic AI & LLM: Agentic AI, Multi-Agent Architectures, LangGraph, CrewAI, AutoGen, OpenAI Agents SDK, MCP, RAG, Agent Memory & State Management

Platform & Client Integration: Kubernetes, Docker, System Design, Distributed Systems, Microservices, REST APIs, React

Cloud & Data: AWS (Lambda, EC2, RDS, SQS, SageMaker), Azure, PostgreSQL, Vector Databases (Pinecone, FAISS), CI/CD

Certifications

AWS Certified Developer Associate | GCP Professional MLOps Engineer