

Ganesh Gupta

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

EXPERIENCE

Oracle | *Software Engineer, AI Platform*

Jun 2025 – Present

- Provisioned and operated **multi-cloud GPU clusters** (AWS, GCP, CoreWeave) with **NVIDIA/CUDA** hardware and **InfiniBand/RoCE** networking for long-running training and **high-availability inference**.
- Wrote systems-level code in **Go, Rust**, and **C++** for **scheduling, resource allocation**, and **fault-tolerant placement** under GPU-scarcity constraints.
- Built observability and **autoscaling** for **vLLM** serving via **Prometheus/Grafana**; ran distributed training via **Slurm** HPC scheduling, keeping utilization up through node loss.
- Built **evaluation harnesses** measuring LLM output quality, safety, **hallucination rates**, and grounding across agent versions; iterated on **eval results** and production usage data.

QureAI | *ML Infrastructure Engineer*

Mar 2026 – Present

- Built **model observability** framework tracking **drift**, latency, and **eval regressions** in production, enabling governance-aligned **rollback** workflows.
- Optimized production LLM inference using **TensorRT-LLM** alongside **vLLM**, tuning serving configuration for **low-latency, high-throughput** generation.

Nomura Research Institute Financial Technologies | *Software Engineer*

Sep 2019 – Aug 2023

- Built and operated **25+ production microservices** in **C++** and SQL-backed **distributed systems** for **1M+ users**, sustaining **99.9% uptime** over 4 years on **AWS** (EC2, RDS, SQS).
- Standardized failure-handling contracts across **25+ financial services**, reducing production escalations **60%** in one quarter via consistent **retry, timeout, and propagation** semantics.
- Standardized API contracts and failure handling across **25+ services**, reducing escalations **60%**; drove **root cause analysis** and **post-incident reviews** to improve long-term reliability.

PROJECTS

Watch-your-LLM

github.com/ganeshhgupta/watch-your-llm

- Production **LLM eval platform** with rubric-based scoring, **drift detection**, and latency tracking across model versions; **FastAPI**, **React/TypeScript**, **PostgreSQL**.

RootCause AI

github.com/ganeshhgupta/rootcause-ai

- LangGraph** multi-agent system diagnosing **production incidents** via tool-use chains; routes alerts to specialized sub-agents and synthesizes **root cause reports**, reducing **MTTR**.

Distributed KV Store / VPC Control Plane Simulator

github.com/ganeshhgupta/distributed-kv-store

- Raft-consensus** KV store with **linearizable reads** and **log compaction**; companion **VPC control plane** simulator modeling subnet routing and security groups.

EDUCATION

University of Texas at Arlington | *M.S. Computer Science, GPA 4.0/4.0*

Aug 2023 – May 2025

MCKV Institute of Engineering | *B.Tech Computer Science*

Jul 2016 – Jul 2020

SKILLS

Languages: Python, Go, Rust, C++, Java, SQL, HTML/CSS

ML Training Infra: PyTorch, LoRA/PEFT Fine-tuning, Distributed Training on Slurm HPC, GPU Utilization, Checkpointing

Cloud and Platform: AWS (SageMaker, EC2, S3), GCP Vertex AI, Kubernetes, Docker, CoreWeave

Reliability: Prometheus, Grafana, Autoscaling, Fault-tolerant Scheduling, On-call

CERTIFICATIONS

AWS Certified Developer Associate | GCP Professional MLOps Engineer