

GANESH GUPTA

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

Experience

Oracle

Jun 2025 – Present

Software Engineer, AI Platform

- Own full-stack delivery across the stack: rich React/TypeScript client experiences powered by GraphQL and RESTful APIs, backed by resilient, high-scale Java/Node.js services at 99%+ uptime.
- Design microservices and data models translating business requirements into technical designs, exposing event-driven REST APIs on multi-tenant AWS (Lambda, SageMaker) and Azure at sub-200ms latency.
- Built a deployment-gated evaluation pipeline enforcing guardrails, security, and observability, personally building the most technically challenging components to block 15% of failed releases.
- Designed stateful recovery via Redis checkpointing for distributed, multithreaded workflows; participate in on-call rotation and lead incident reviews.
- Established engineering standards for CI/CD, testing, and security across the platform, mentoring engineers through architecture reviews and design write-ups.

QureAI (AI Startup)

Mar 2026 – Present

ML Infrastructure Engineer

- Integrated AWS Bedrock-hosted foundation models, including Claude, into production agentic AI workflows for managed, scalable model access as an early engineer at this AI startup.
- Improved RAG retrieval recall 61% to 89% on a 2K-query eval set by rebuilding the chunking pipeline with hybrid dense-sparse retrieval and re-ranking.

Nomura Research Institute Financial Technologies

Sep 2019 – Aug 2023

Software Engineer

- Architected and operated 25+ production Java Spring Boot microservices exposing REST APIs for 1M+ users, sustaining 99.9% uptime over 4 years on AWS (EC2, RDS, SQS).
- Drove a redesign of nightly data aggregation, cutting processing from 8 minutes to 90 seconds via indexed views, T-SQL optimization, and batch partitioning.
- Standardized failure-handling contracts (retry, timeout, propagation semantics) across 25+ services, reducing production escalations 60% in one quarter.
- Mentored engineers through code reviews and design write-ups, driving consistent engineering standards across the team.

Projects

LangFetch | tinyurl.com/LangFetch

- LangGraph agent compiling natural-language questions into dependency-ordered SQL subqueries with Pinecone vector retrieval and hybrid re-ranking; deployed as an MCP-compatible server for enterprise data queries.

RootCause AI | github.com/ganeshhgupta/rootcause-ai

- LangGraph multi-agent system diagnosing production incidents via tool-use chains; routes alerts to specialized sub-agents and synthesizes evidence-backed root cause reports, reducing mean time to resolution.

Watch-your-LLM | github.com/ganeshhgupta/watch-your-llm

- Production LLM eval platform with hallucination detection, drift tracking, rubric-based scoring, and regression alerting across model versions; FastAPI backend, PostgreSQL.

Education

University of Texas at Arlington

Aug 2023 – May 2025

MS, Computer Science, GPA 4.0/4.0

MCKV Institute of Engineering

Jul 2016 – Jul 2020

BTech, Computer Science

Skills

Languages: Java, TypeScript, Python, SQL, C++

Frontend: React, TypeScript, HTML/CSS, Component Architecture, State Management

Backend & Security: Java, Spring Boot, Node.js, REST APIs, GraphQL, Microservices, OAuth, RBAC

Data & Cloud: SQL, NoSQL (MongoDB, DynamoDB), AWS, Azure, Docker, CI/CD

Certifications

AWS Certified Developer Associate | GCP Professional MLOps Engineer