

GANESH GUPTA

817-822-1131 | iamgs10rk@gmail.com | linkedin.com/in/ganeshhgupta | github.com/ganeshhgupta

Experience

Oracle

Jun 2025 – Present

Software Engineer, AI Platform

- Own multi-step agent workflows end-to-end using LangGraph, CrewAI, AutoGen, and OpenAI Agents SDK, with tool-calling integrations and stateful agent memory via Redis checkpointing at 99%+ uptime.
- Combined GraphRAG (knowledge-graph-augmented retrieval) with vector-based RAG pipelines into a unified retrieval strategy, improving grounding accuracy on complex, multi-hop enterprise queries.
- Built a deployment-gated evaluation framework enforcing guardrails, AI governance, and observability around agent outputs, blocking 15% of failed releases before they reached production.
- Build full-stack AI features (React/TypeScript frontend, Python backend) partnering directly with product and business stakeholders to turn real problems into deployed AI solutions.
- Set engineering standards for CI/CD, testing, and secure development practices across the platform, working in Agile sprints with Git-based version control.

QureAI (AI Startup)

Mar 2026 – Present

ML Infrastructure Engineer

- Built custom evaluators for RAG/agent workflows tracking drift, latency, and regression across model versions, catching failures before they reached production.
- Improved semantic search/RAG retrieval recall 61% to 89% on a 2K-query eval set by rebuilding the chunking pipeline with hybrid dense-sparse retrieval and re-ranking.

Nomura Research Institute Financial Technologies

Sep 2019 – Aug 2023

Software Engineer

- Architected and operated 25+ production C++ and SQL-backed microservices for 1M+ users, sustaining 99.9% uptime over 4 years on AWS (EC2, RDS, SQS).
- Drove a redesign of nightly data aggregation, cutting processing from 8 minutes to 90 seconds via indexed views and batch partitioning.
- Standardized failure-handling contracts (retry, timeout, propagation semantics) across 25+ services, reducing production escalations 60% in one quarter.
- Mentored engineers through code reviews and design write-ups, driving consistent engineering standards.

Projects

LangFetch | tinyurl.com/LangFetch

- LangGraph agent compiling natural-language questions into dependency-ordered SQL subqueries with Pinecone vector retrieval and hybrid re-ranking; deployed as an MCP-compatible server for enterprise data queries.

RootCause AI | github.com/ganeshhgupta/rootcause-ai

- LangGraph multi-agent system diagnosing production incidents via tool-use chains; routes alerts to specialized sub-agents and synthesizes evidence-backed root cause reports, reducing mean time to resolution.

Watch-your-LLM | github.com/ganeshhgupta/watch-your-llm

- Production LLM eval platform with hallucination detection, drift tracking, rubric-based scoring, and regression alerting across model versions; FastAPI backend, PostgreSQL.

Education

University of Texas at Arlington

Aug 2023 – May 2025

MS, Computer Science, GPA 4.0/4.0

MCKV Institute of Engineering

Jul 2016 – Jul 2020

BTech, Computer Science

Skills

Languages: Python, TypeScript, SQL, Java, C++

Agentic AI & Orchestration: LangGraph, CrewAI, AutoGen, OpenAI Agents SDK, Anthropic Claude, MCP, Tool Calling, Agent Memory/State

Retrieval & Evaluation: RAG, GraphRAG, Custom Evaluators, Model/Agent Observability, Guardrails

Full-Stack & Cloud: React, REST APIs, Vector DBs (Pinecone, FAISS), AWS, Azure, GCP, CI/CD

Certifications

AWS Certified Developer Associate | GCP Professional MLOps Engineer